

Using OpenCV for Science



https://github.com/Aramist/opencv_tutorial

2024-09-17

Table of Contents

1. What is OpenCV
2. Installing OpenCV
3. Data structures
4. Common Image Operations
5. Hands on: Contours & Basic Object Detection
6. Uncommon Image Operations
7. Working with GUIs
8. Case study: Data Annotation
9. Working with Videos
10. Challenge: Football Player Tracking

What is OpenCV?

- A very expansive C++ imaging processing library with a Python wrapper

- Main modules:

- core. **Core functionality**
- imgproc. **Image Processing**
- imgcodecs. **Image file reading and writing**
- videoio. **Video I/O**
- highgui. **High-level GUI**
- video. **Video Analysis**
- calib3d. **Camera Calibration and 3D Reconstruction**
- features2d. **2D Features Framework**
- objdetect. **Object Detection**
- dnn. **Deep Neural Network module**
- ml. **Machine Learning**
- flann. **Clustering and Search in Multi-Dimensional Spaces**
- photo. **Computational Photography**
- stitching. **Images stitching**
- gapi. **Graph API**

- Extra modules:

- alphamat. **Alpha Matting**
- aruco. **Aruco markers**, module functionality was moved to objdetect module
- bgsegm. **Improved Background-Foreground Segmentation Methods**
- bioinspired. **Biologically inspired vision models and derived tools**
- cannops. **Ascend-accelerated Computer Vision**
- ccalib. **Custom Calibration Pattern for 3D reconstruction**
- cudaarithm. **Operations on Matrices**
- cudabgsegm. **Background Segmentation**
- cudacodec. **Video Encoding/Decoding**
- cudafeatures2d. **Feature Detection and Description**
- cudafilters. **Image Filtering**
- cudaimproc. **Image Processing**
- cudalegacy. **Legacy support**
- cudaobjdetect. **Object Detection**
- cudaoptflow. **Optical Flow**
- cudastereo. **Stereo Correspondence**
- cudawarping. **Image Warping**
- cudev. **Device layer**
- cvv. **GUI for Interactive Visual Debugging of Computer Vision Programs**
- datasets. **Framework for working with different datasets**
- dnn_objdetect. **DNN used for object detection**
- dnn_superres. **DNN used for super resolution**
- dpm. **Deformable Part-based Models**
- face. **Face Analysis**
- freetype. **Drawing UTF-8 strings with freetype/harfbuzz**
- fuzzy. **Image processing based on fuzzy mathematics**
- hdf. **Hierarchical Data Format I/O routines**
- hfs. **Hierarchical Feature Selection for Efficient Image Segmentation**
- img_hash. **The module brings implementations of different image hashing algorithms.**
- intensity_transform. **The module brings implementations of intensity transformation algorithms.**
- julia. **Julia bindings for OpenCV**
- line_descriptor. **Binary descriptors for lines extracted from an image**
- mcc. **Macbeth Chart module**
- optflow. **Optical Flow Algorithms**
- ovis. **OGRE 3D Visualiser**
- phase_unwrapping. **Phase Unwrapping API**
- plot. **Plot function for Mat data**
- quality. **Image Quality Analysis (IQA) API**
- rapid. **silhouette based 3D object tracking**
- reg. **Image Registration**
- rgbd. **RGB-Depth Processing**
- saliency. **Saliency API**
- sfm. **Structure From Motion**
- shape. **Shape Distance and Matching**
- signal. **Signal Processing**
- stereo. **Stereo Correspondance Algorithms**
- structured_light. **Structured Light API**
- superres. **Super Resolution**
- surface_matching. **Surface Matching**
- text. **Scene Text Detection and Recognition**
- tracking. **Tracking API**
- videotab. **Video Stabilization**
- viz. **3D Visualizer**
- wechat_qrcode. **WeChat QR code detector for detecting and parsing QR code.**
- xfeatures2d. **Extra 2D Features Framework**
- ximgproc. **Extended Image Processing**
- xobjdetect. **Extended object detection**
- xphoto. **Additional photo processing algorithms**

- A very expansive C++ imaging processing library with a Python wrapper
- Error messages are usually aren't very informative

```
>>> H = cv2.getPerspectiveTransform(src_points[...], None, None], dst_points)
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
cv2.error: OpenCV(4.8.0) /Users/xperience/GHA-OpenCV-Python/_work/opencv-python/
opencv-python/opencv/modules/imgproc/src/imgwarp.cpp:3528: error: (-215:Assertion failed) src.checkVector(2, CV_32F) == 4 && dst.checkVector(2, CV_32F) == 4 in
function 'getPerspectiveTransform'
```

- Documentation is C++ oriented with Python on the sidelines

• `getPerspectiveTransform()` [2/2]

C++:

```
Mat cv::getPerspectiveTransform ( InputArray src,
                                 InputArray dst,
                                 int          solveMethod = DECOMP_LU
                               )
```

Python:

```
cv.getPerspectiveTransform( src, dst[, solveMethod] ) -> retval
```

Installing OpenCV

`pip install opencv-python`

Data Structures

Data Structures

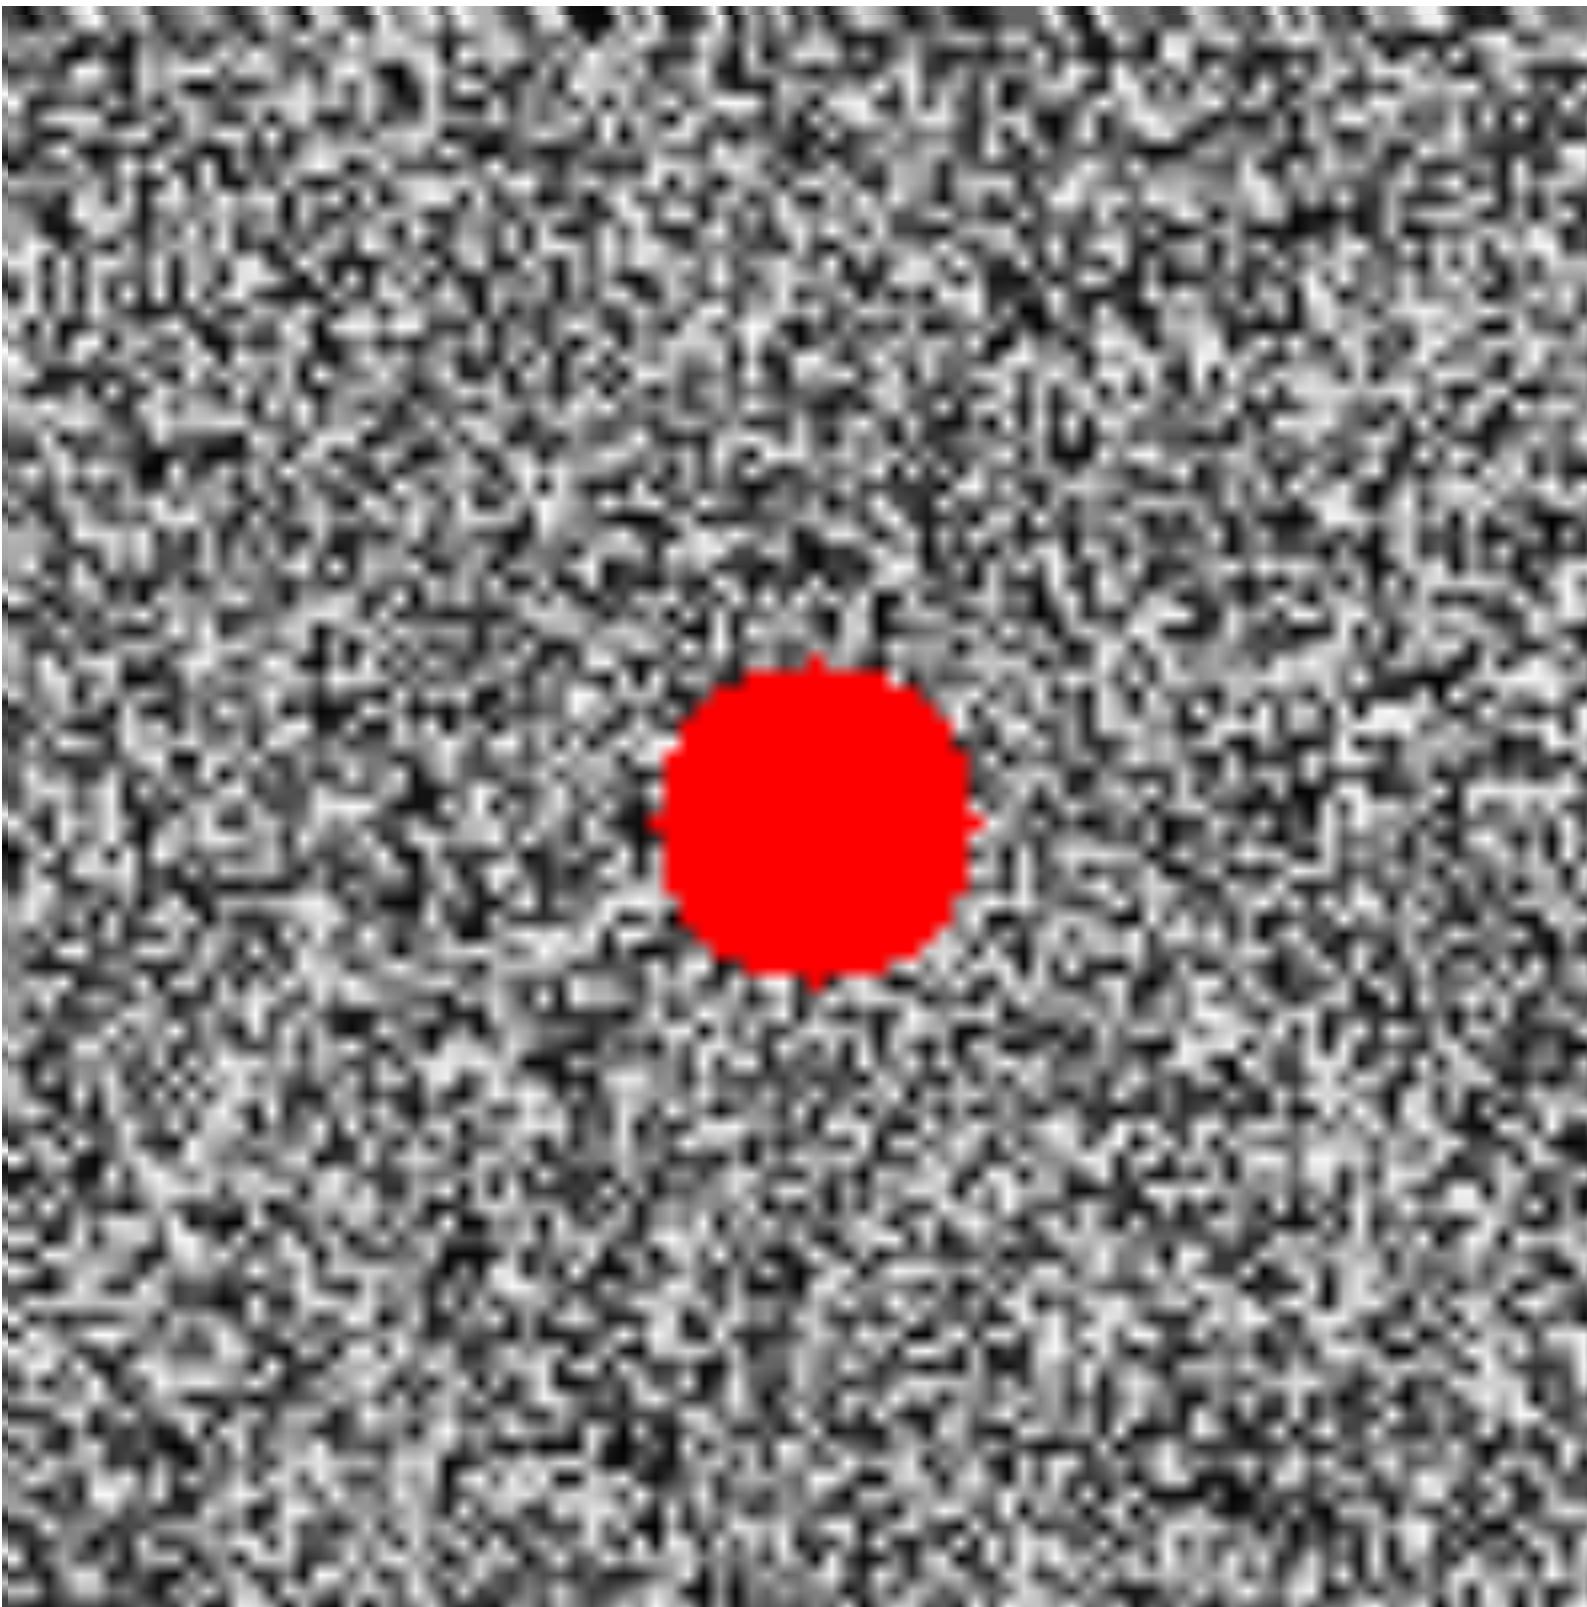
- One of the more confusing things to OpenCV beginners
 - Images are NumPy arrays
 - C-ordering (row index first, then column)
 - $\text{img}[y, x]$ -> pixel at (x, y)
 - $Y=0$ at top of image (like a matrix in mathematics)
 - When manipulating images directly, use (y, x) order (e.g. $\text{img}[y, x, :] = (0, 0, 255)$)
 - When calling OpenCV functions, use (x, y) order (e.g. `cv2.circle(img, (x, y), ...)`)

Data Structures

- Grayscale images have shape (height, width) and dtype np.uint8
- Colored images have shape (height, width, num_channels) and (usually) dtype np.uint8. Some functions support greater bit depths
- Non-image matrices are usually CV_32F (OpenCV's internal name for float32 arrays)
 - Will show up if you are fitting rectangles, finding contours, or transforming images

Common Image Operations

Common Image Operations: Color Space Conversion



```
graycircle = cv2.cvtColor(circle, cv2.COLOR_BGR2GRAY)
```

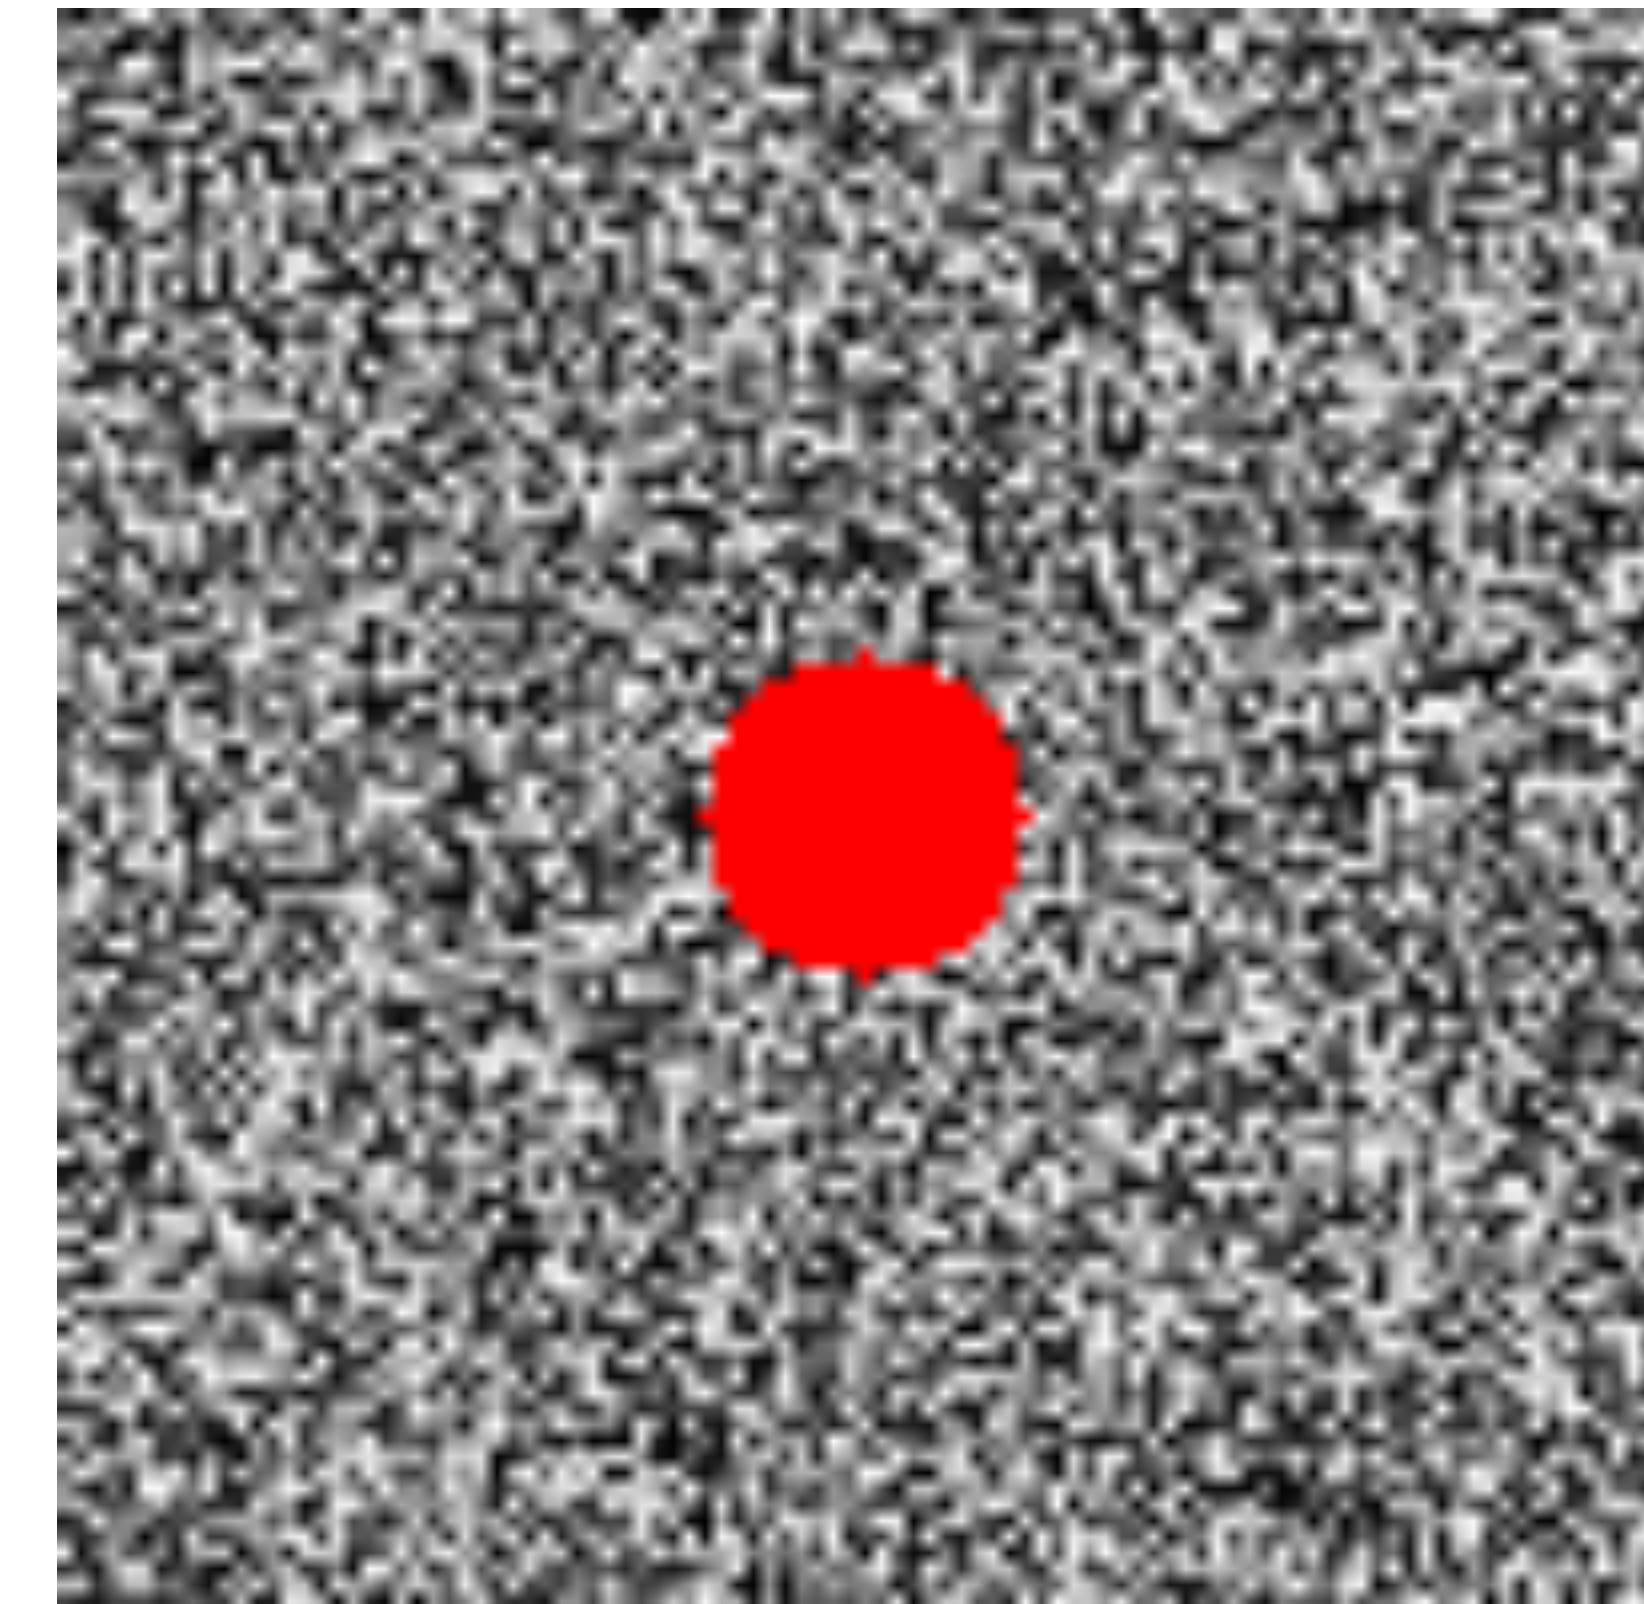
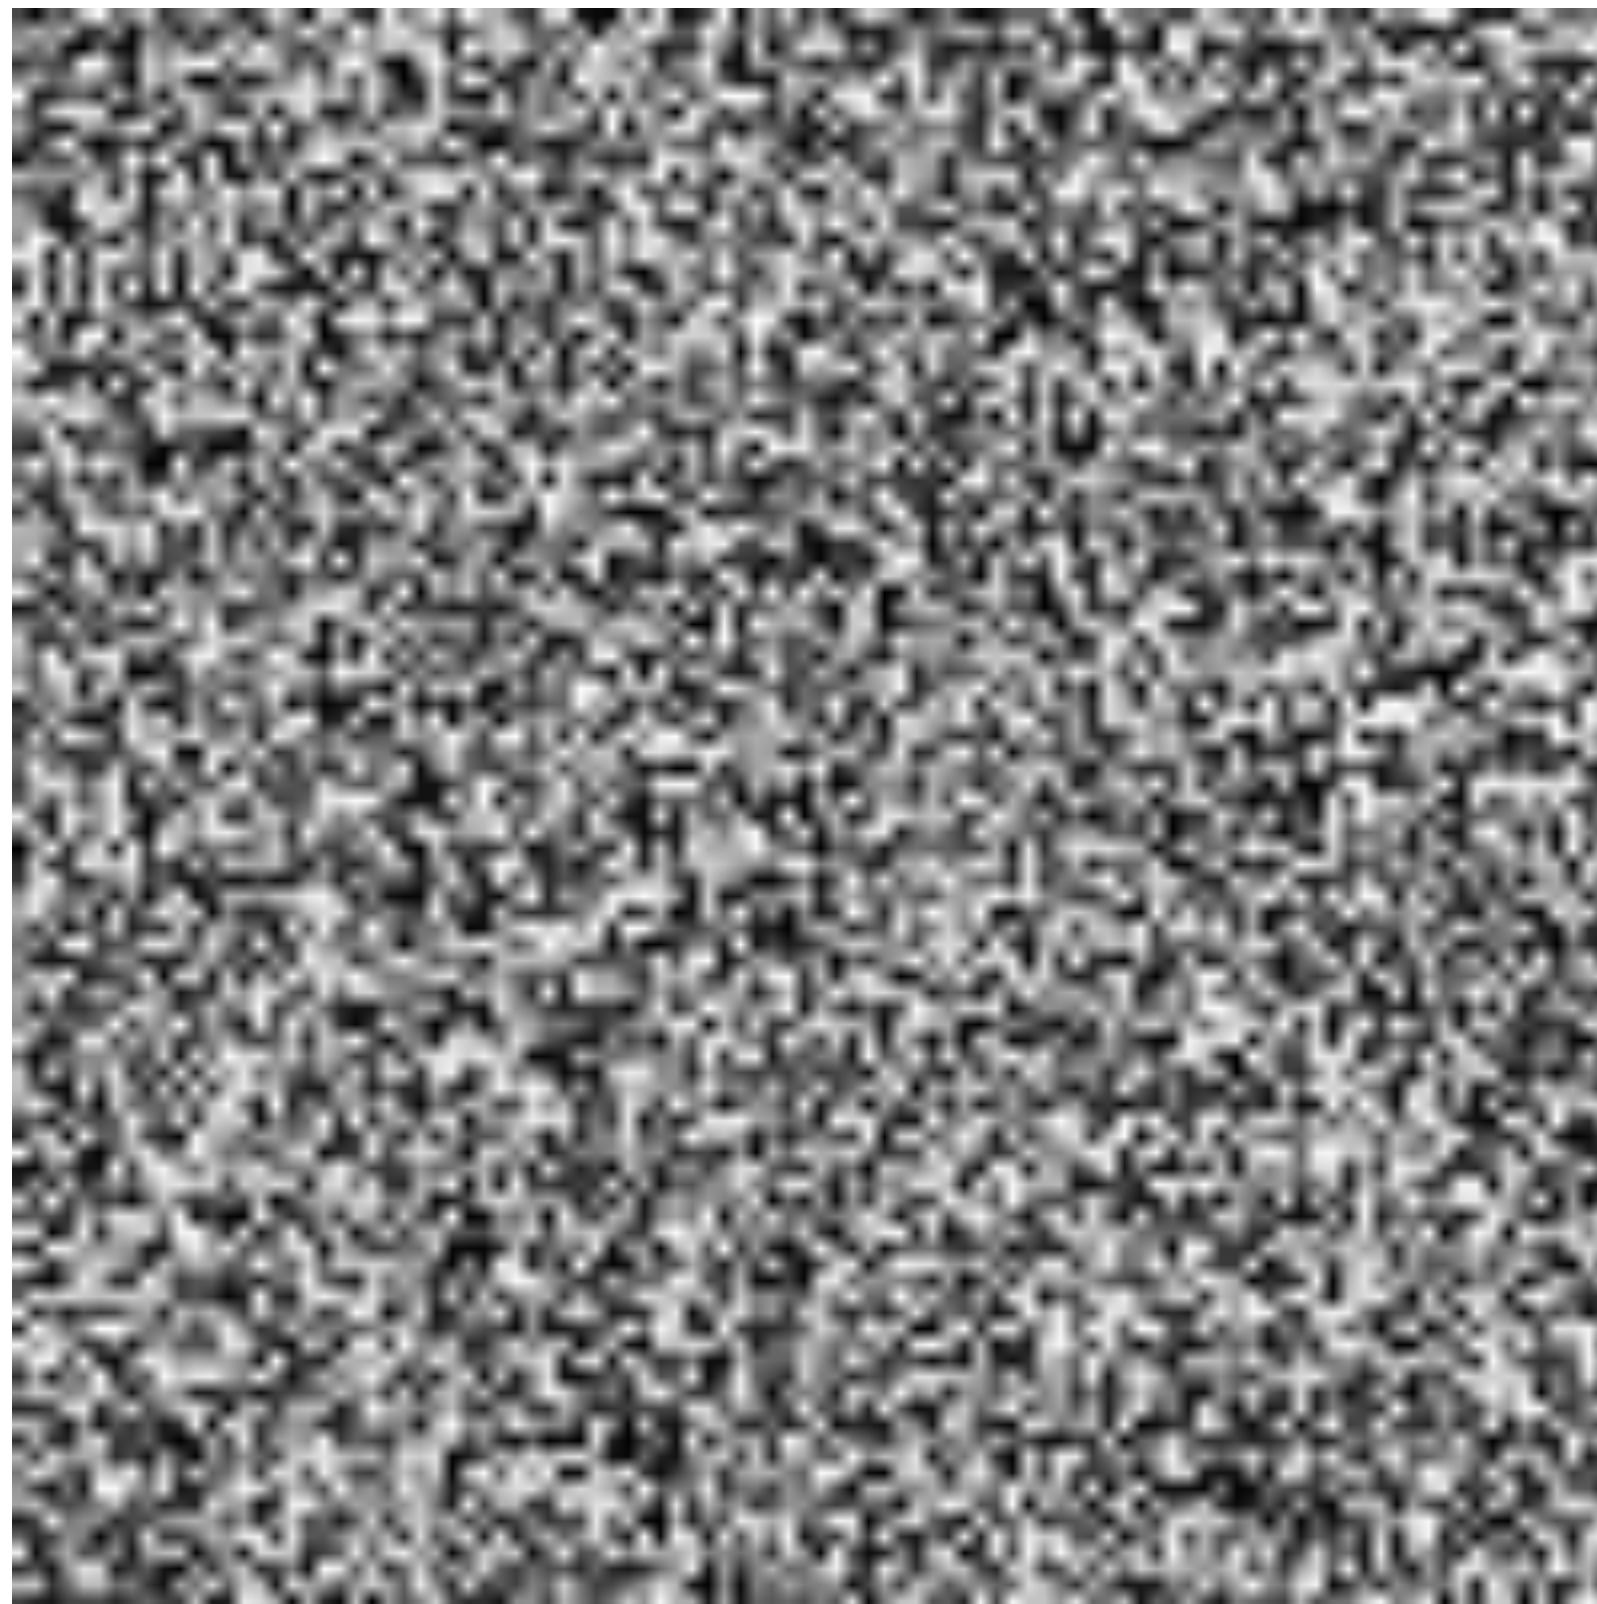
Common Image Operations: Color Space Conversion

- Arguably the most useful function in object detection
 - "A lot of problems are easily solved if you choose the right basis" – Eero, probably
- Supports every color space you know, and then some
- The only (multichannel) OpenCV functions which make any assumptions about the image's color space are imread and imwrite, which assume BGR or grayscale

Common Image Operations: Drawing

<- 100px ->

^ - 100px ->
v

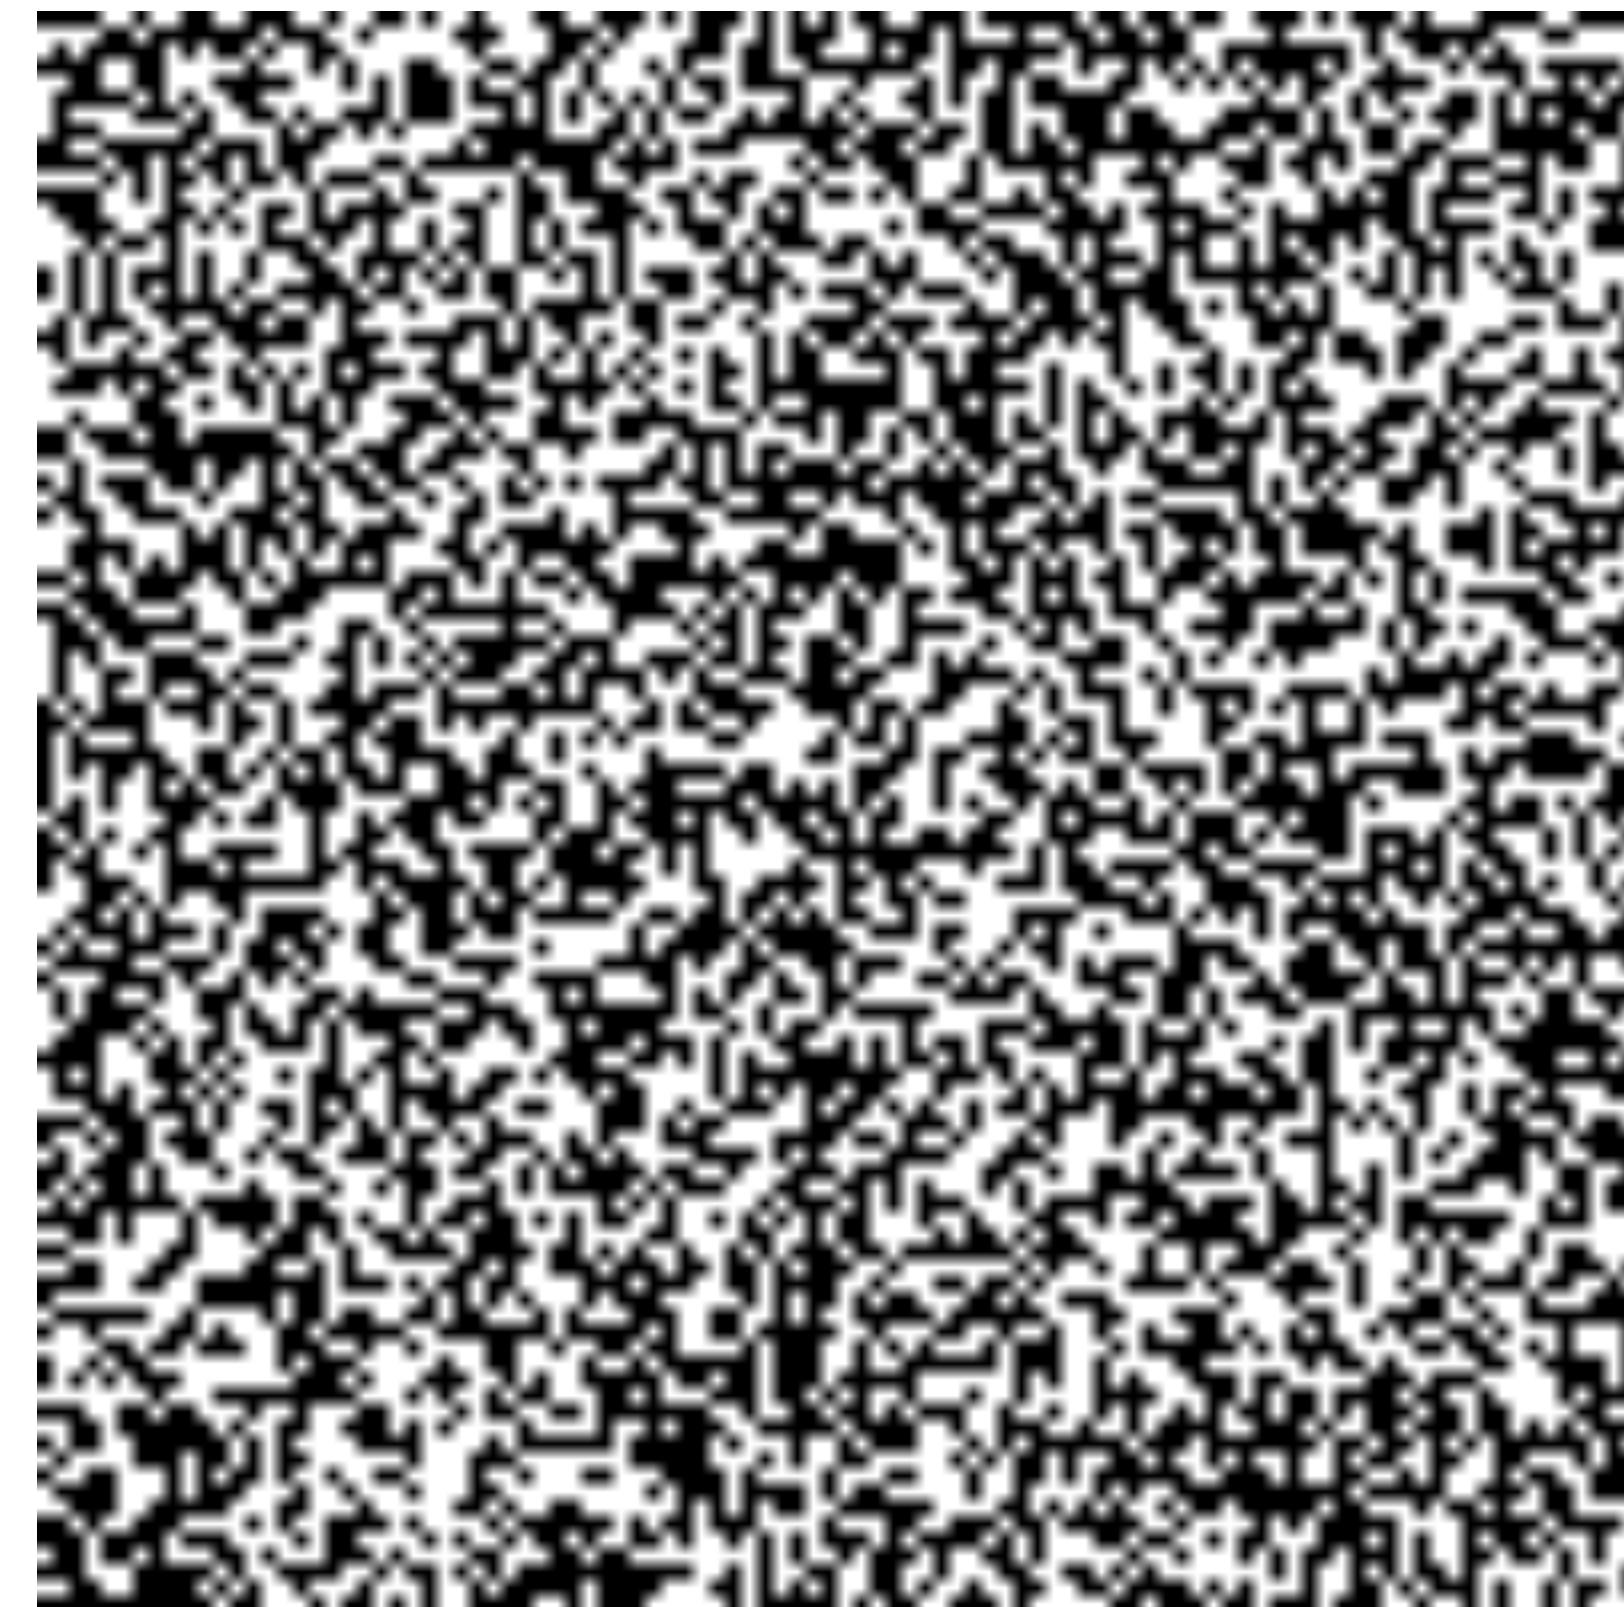
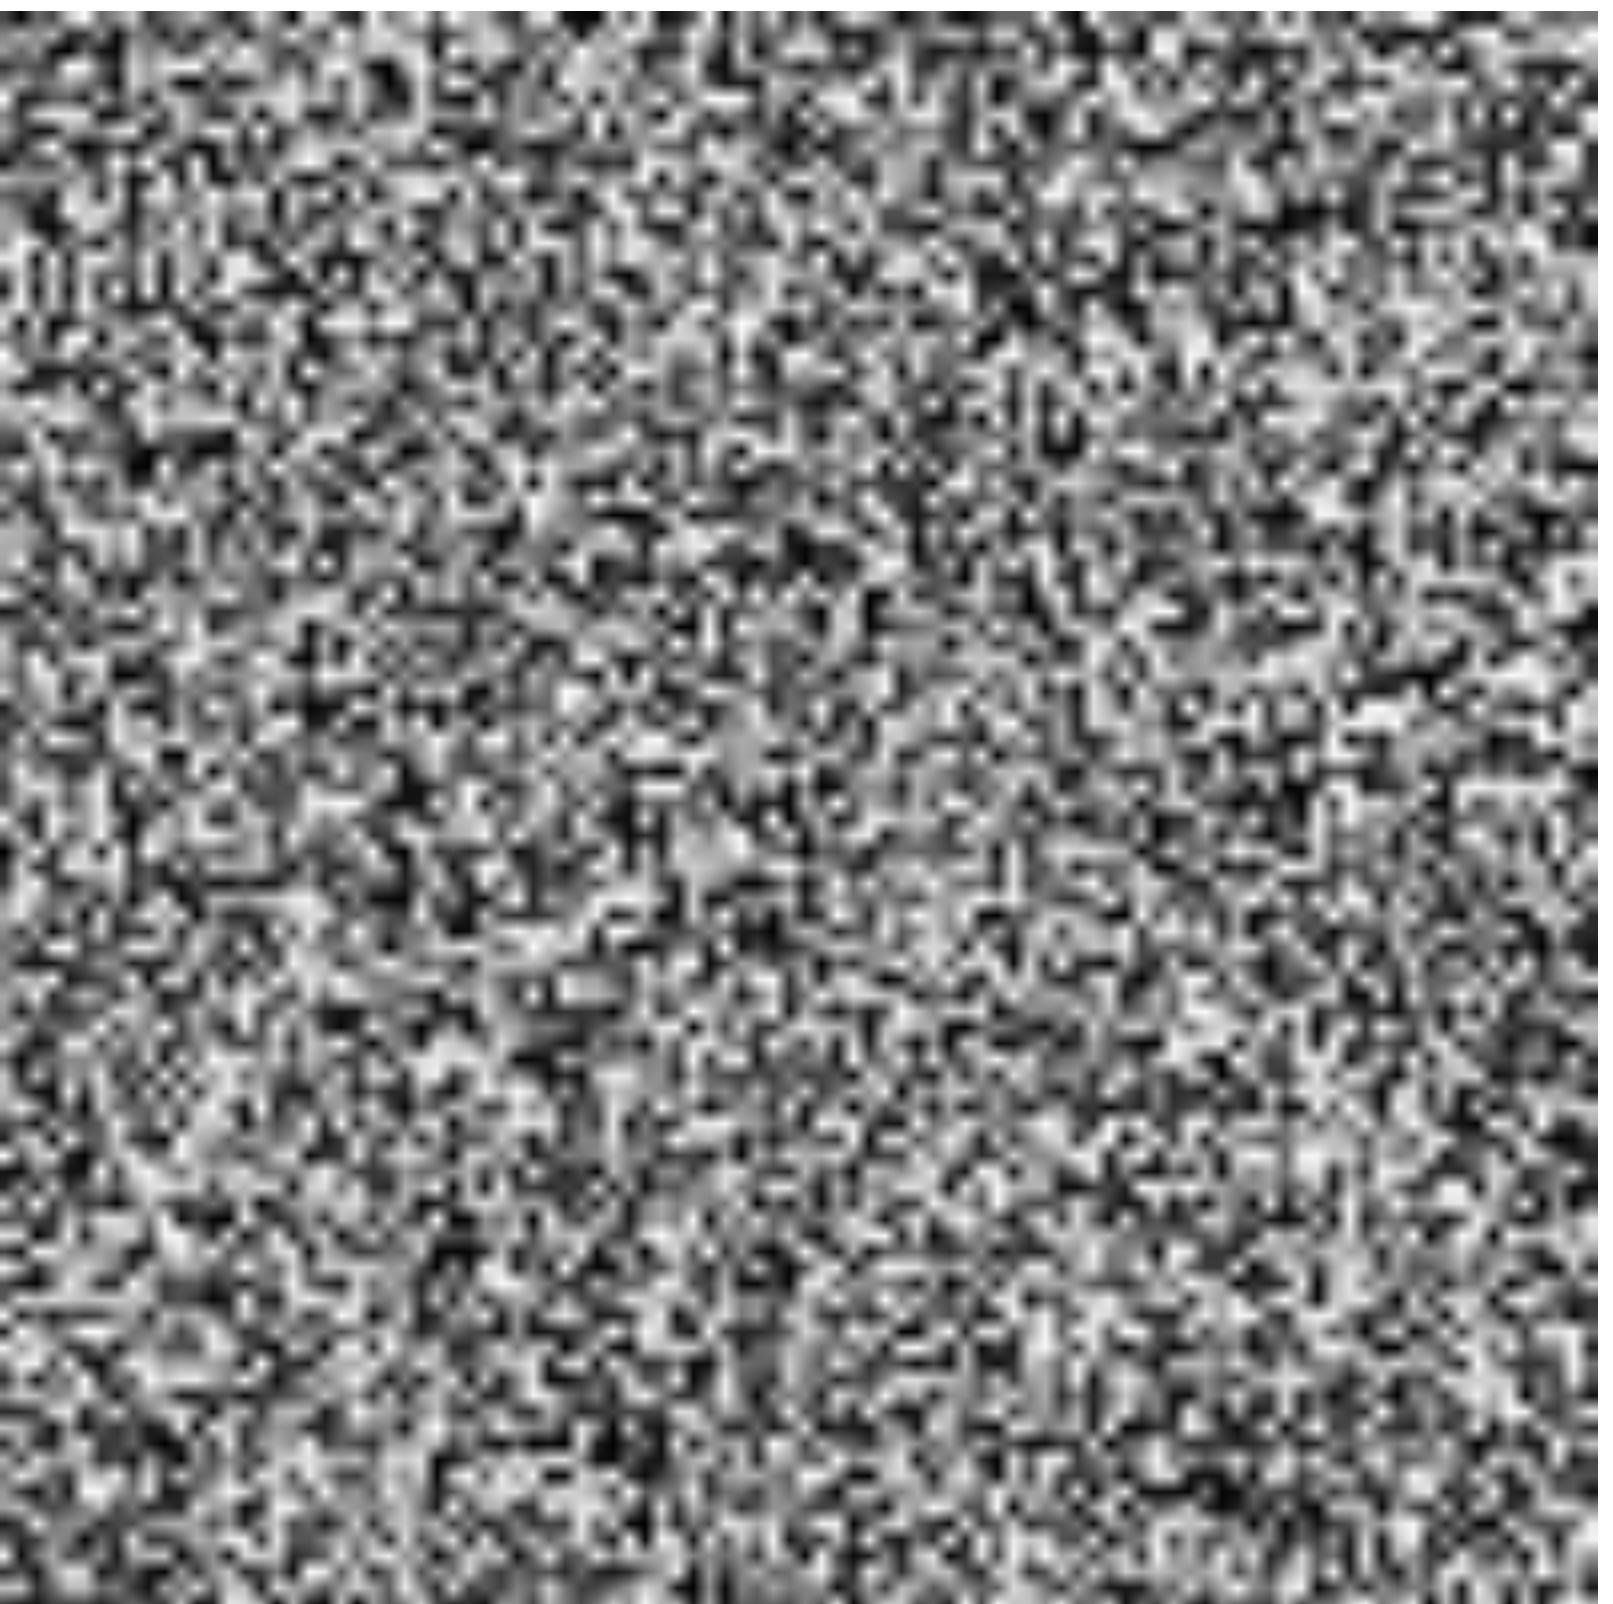


```
im_w_circle = cv2.circle(noise, (50, 50), 10, (0, 0, 255), -1)
```

Common Image Operations: Drawing

- Also see: cv2.drawContours, cv2.putText, cv2.ellipse, cv2.line, cv2.rectangle

Common Image Operations: Thresholding

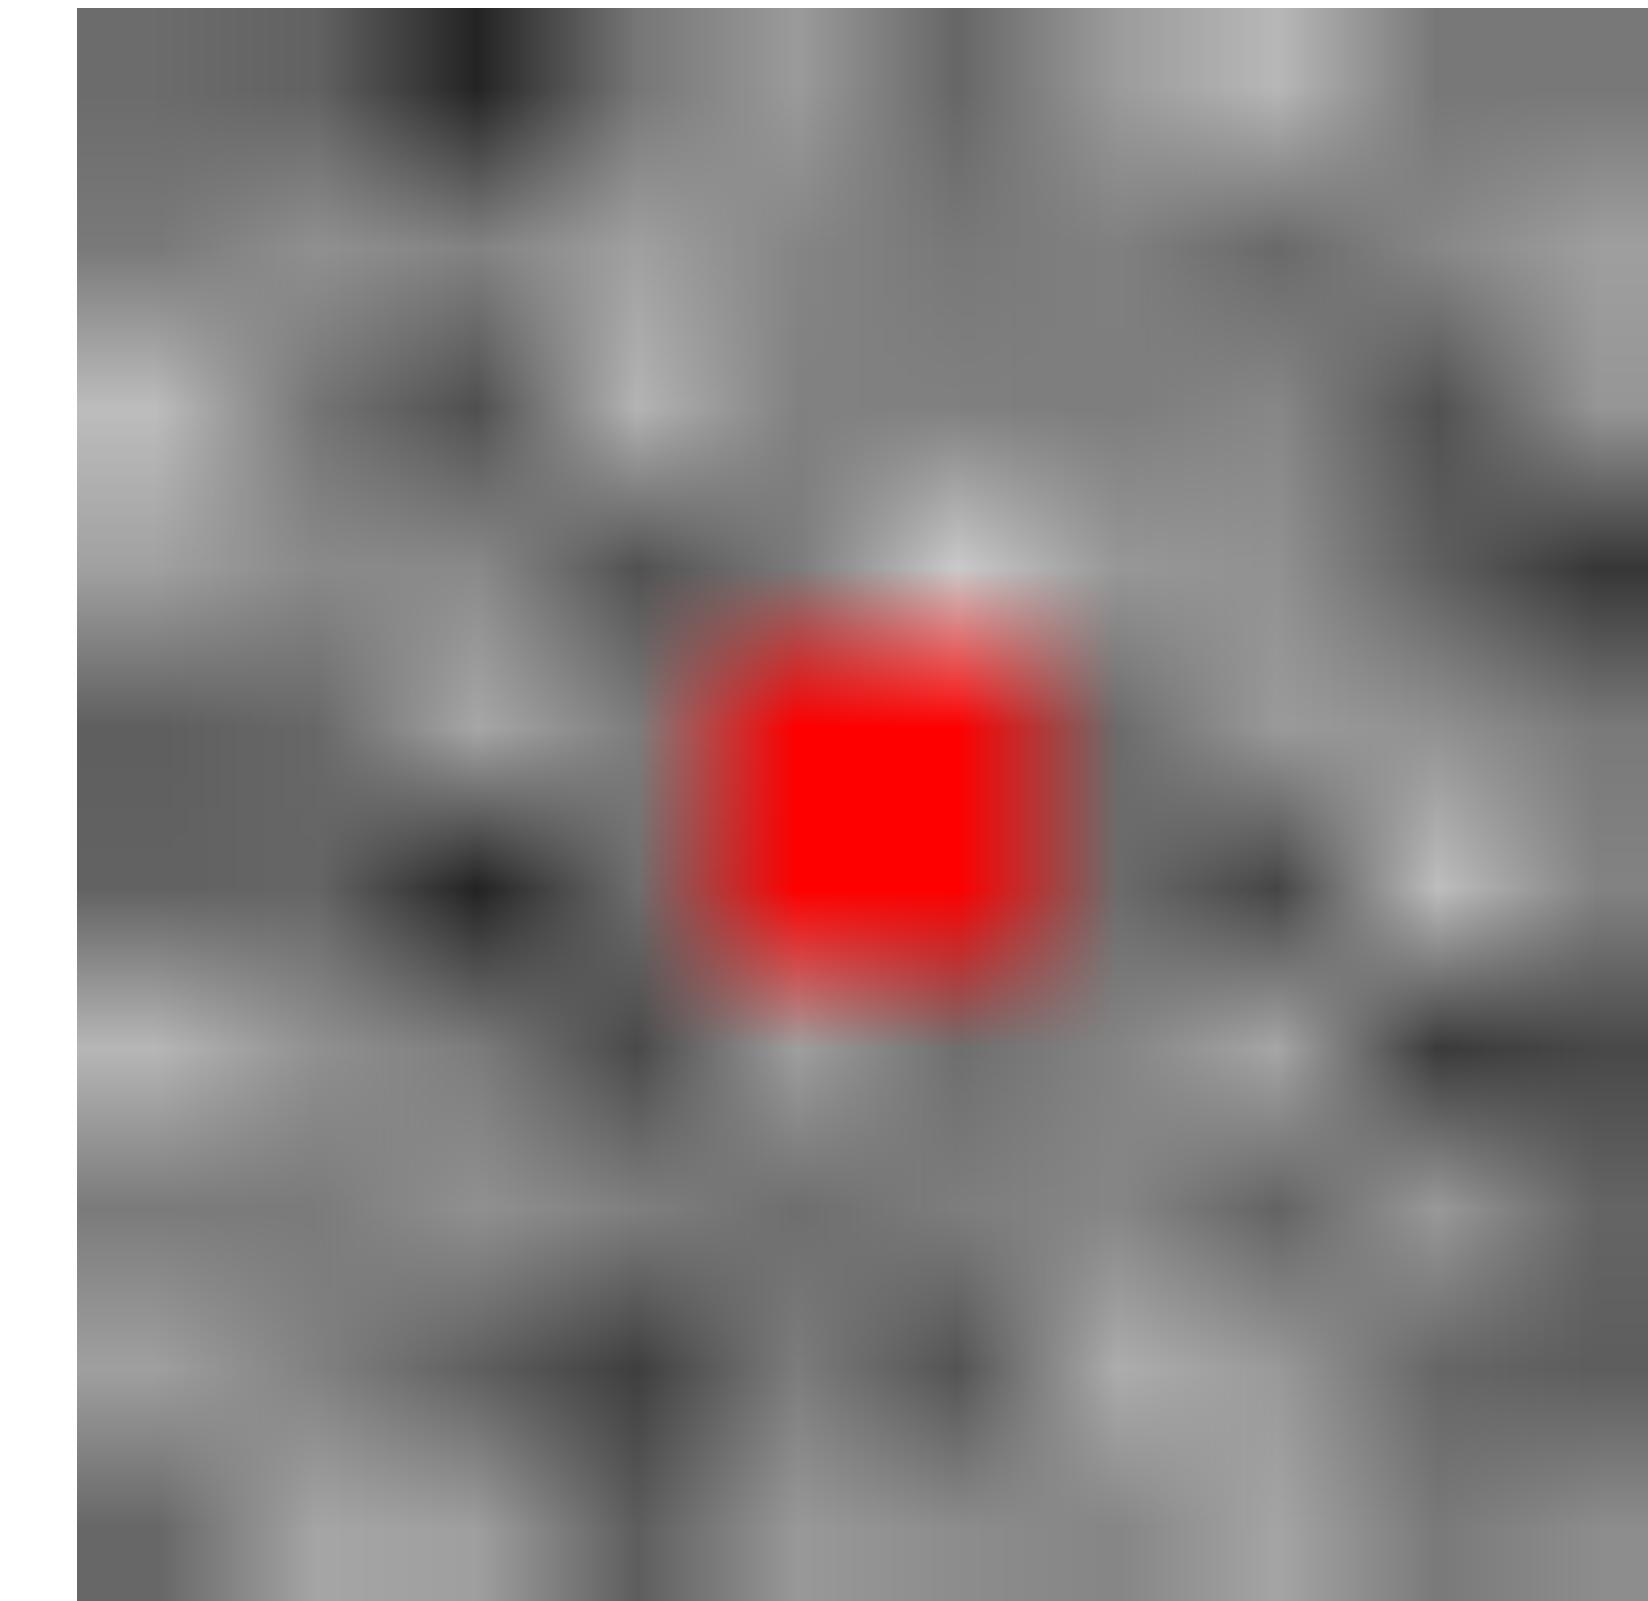
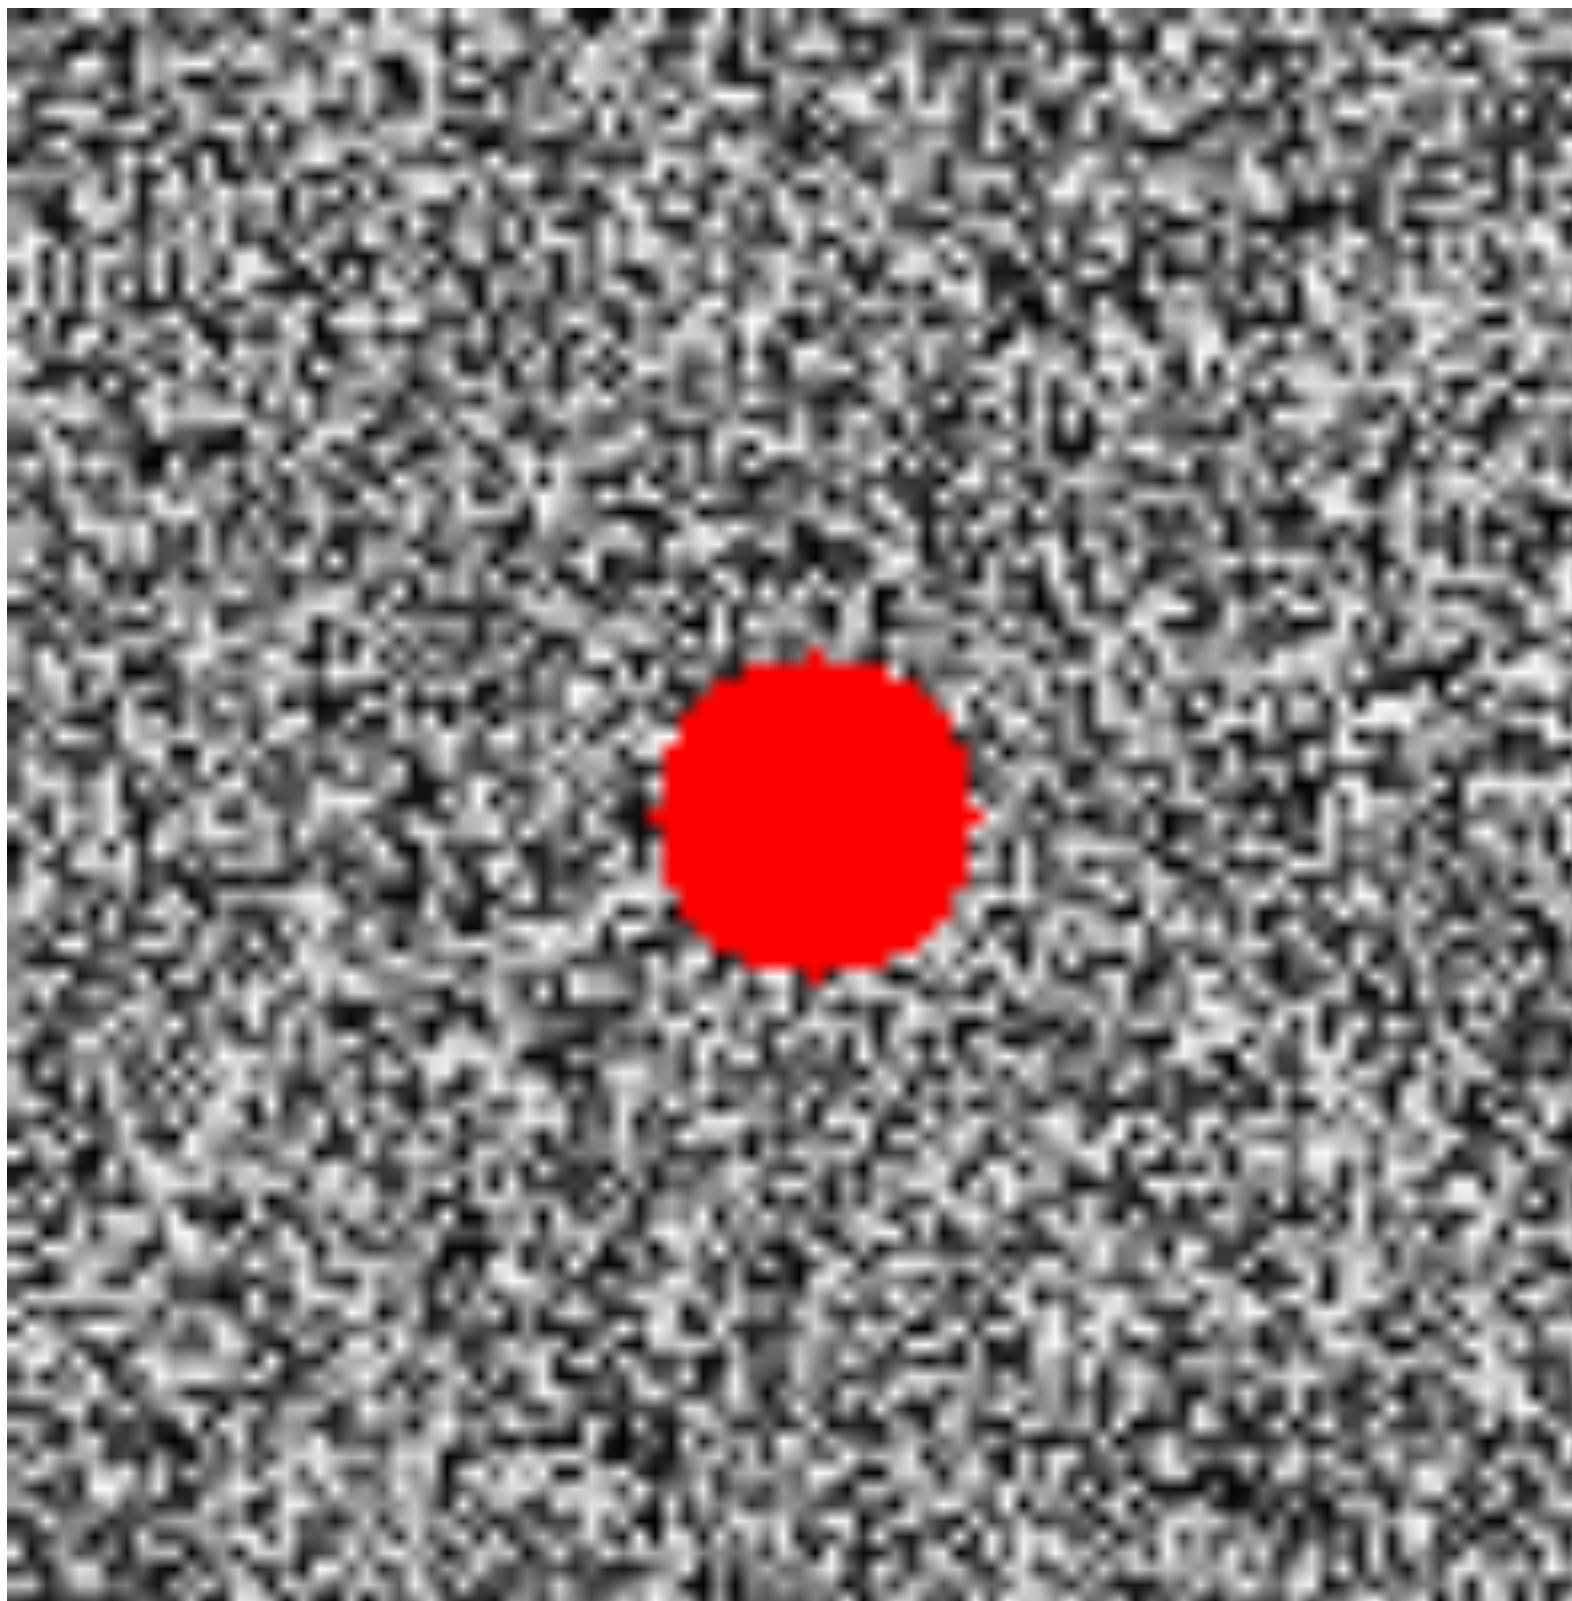


```
_, thresh = cv2.threshold(im, 127, 255, cv2.THRESH_BINARY)
```

Common Image Operations: Thresholding

- Good for masking images
- Has many modes (see: [`cv::ThresholdTypes`](#))
- In practice I usually just use NumPy unless I'm doing adaptive thresholding (see: [`cv2.adaptiveThreshold`](#))
 - The previous operation can be rewritten in many ways:
 - `Im = np.where(im > 127, 255, 0).astype(np.uint8)`
 - `Im = (im > 127).astype(np.uint8) * 255`
 - ...

Common Image Operations: Resizing

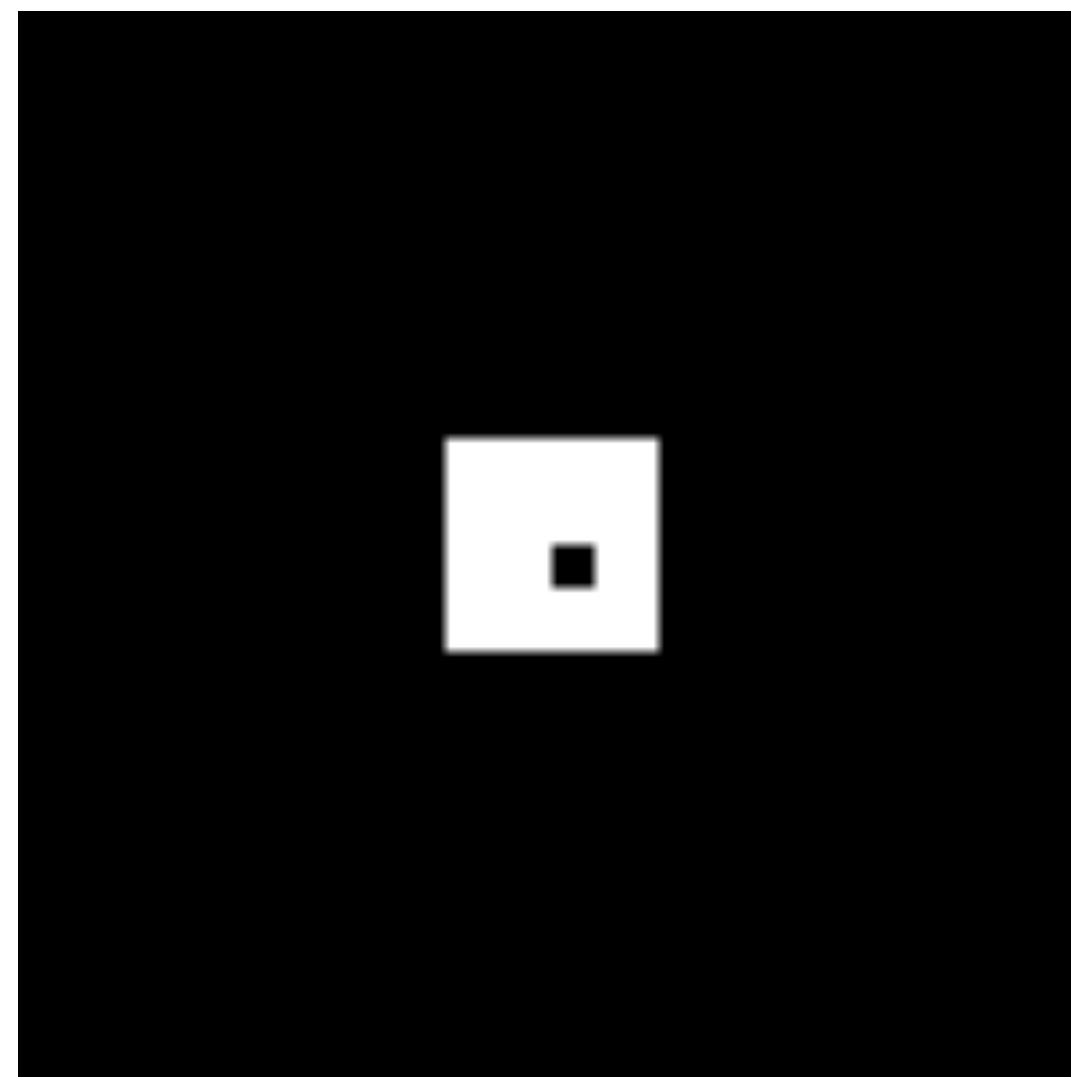


```
circle_small = cv2.resize(circle, (10, 10))
```

Common Image Operations: Resizing

- Not much to say
- Supports several interpolation methods, default is linear

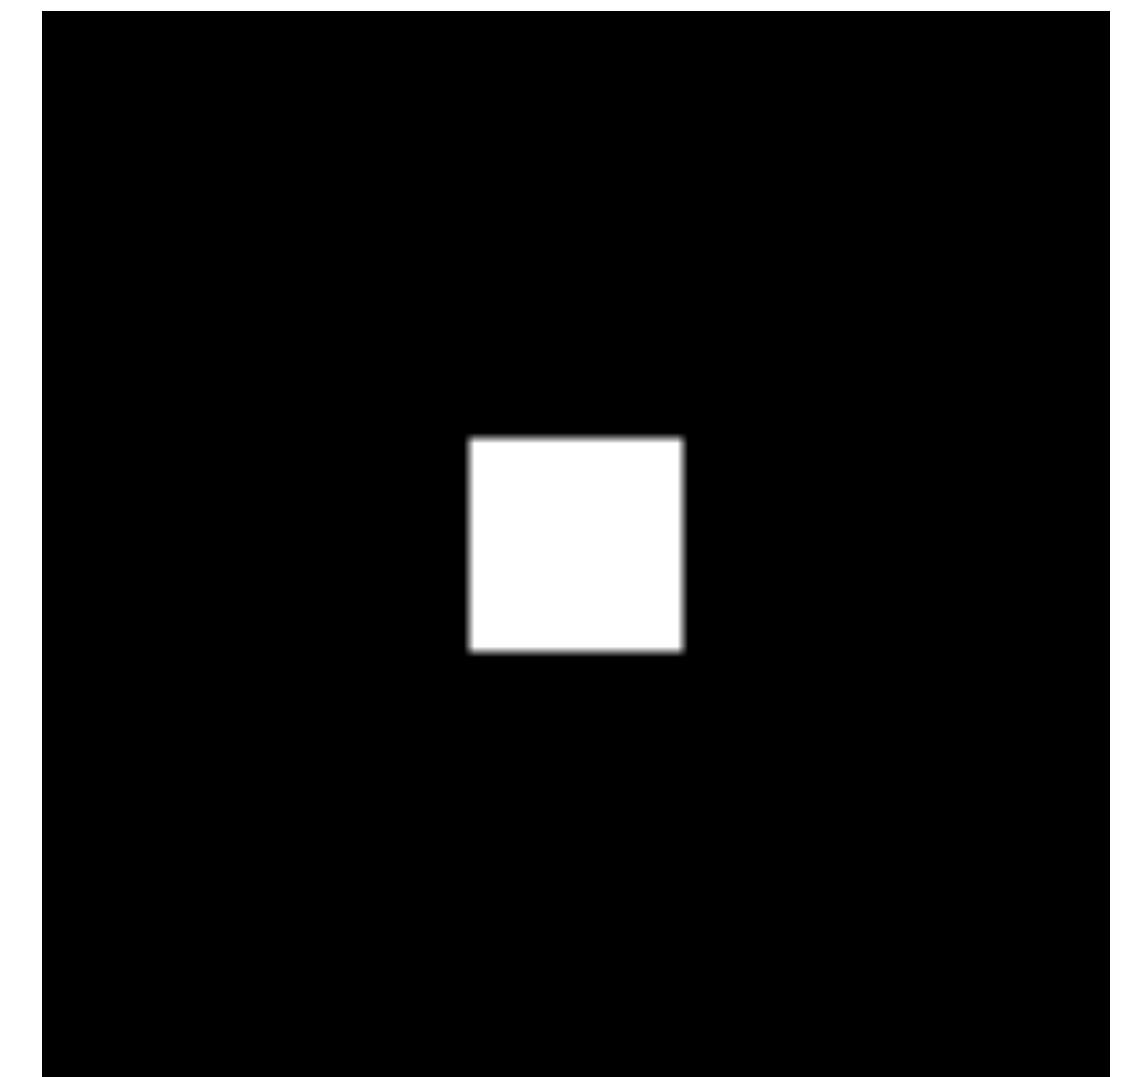
Common Image Operations: Erode & Dilate



-> Dilate ->



-> Erode ->



```
kernel = np.ones((5, 5), dtype=np.uint8)
dilation = cv2.dilate(im_with_hole, kernel)
eroded = cv2.erode(dilation, kernel)
```

Common Image Operations: Erode & Dilate

- Used all the time in object detection to:
 - Patch holes within solid objects (dilate then erode)
 - Remove small artifacts while preserving large objects (erode then dilate)
- The structuring element (improperly labeled kernel in the last slide) determines the "rule" for which pixels get turned on or off

Honorable Mentions:

- Blur: [cv2.GaussianBlur](#), [cv2.blur](#)
- Save your arrays: [np.ascontiguousarray](#)

Hands on: Basic Object Detection

Hands on: Contours & Basic Object Detection

- Detect contours in binary images: `cv2.findContours(image, mode, approx_method)`
- Mode determines how OpenCV handles the contour hierarchy. For basic object detection, `cv2.RETR_LIST` or `cv2.RETR_EXTERNAL` are fine. If you care about whether a contour lies within another contour, consider reading the wiki (see: [cv::RetrievalModes](#))
- `approx_method` can usually be ignored, determines how contour edges are approximated

Hands on: Contours & Basic Object Detection

How can we find
the mouse?



Hands on: Contours & Basic Object Detection



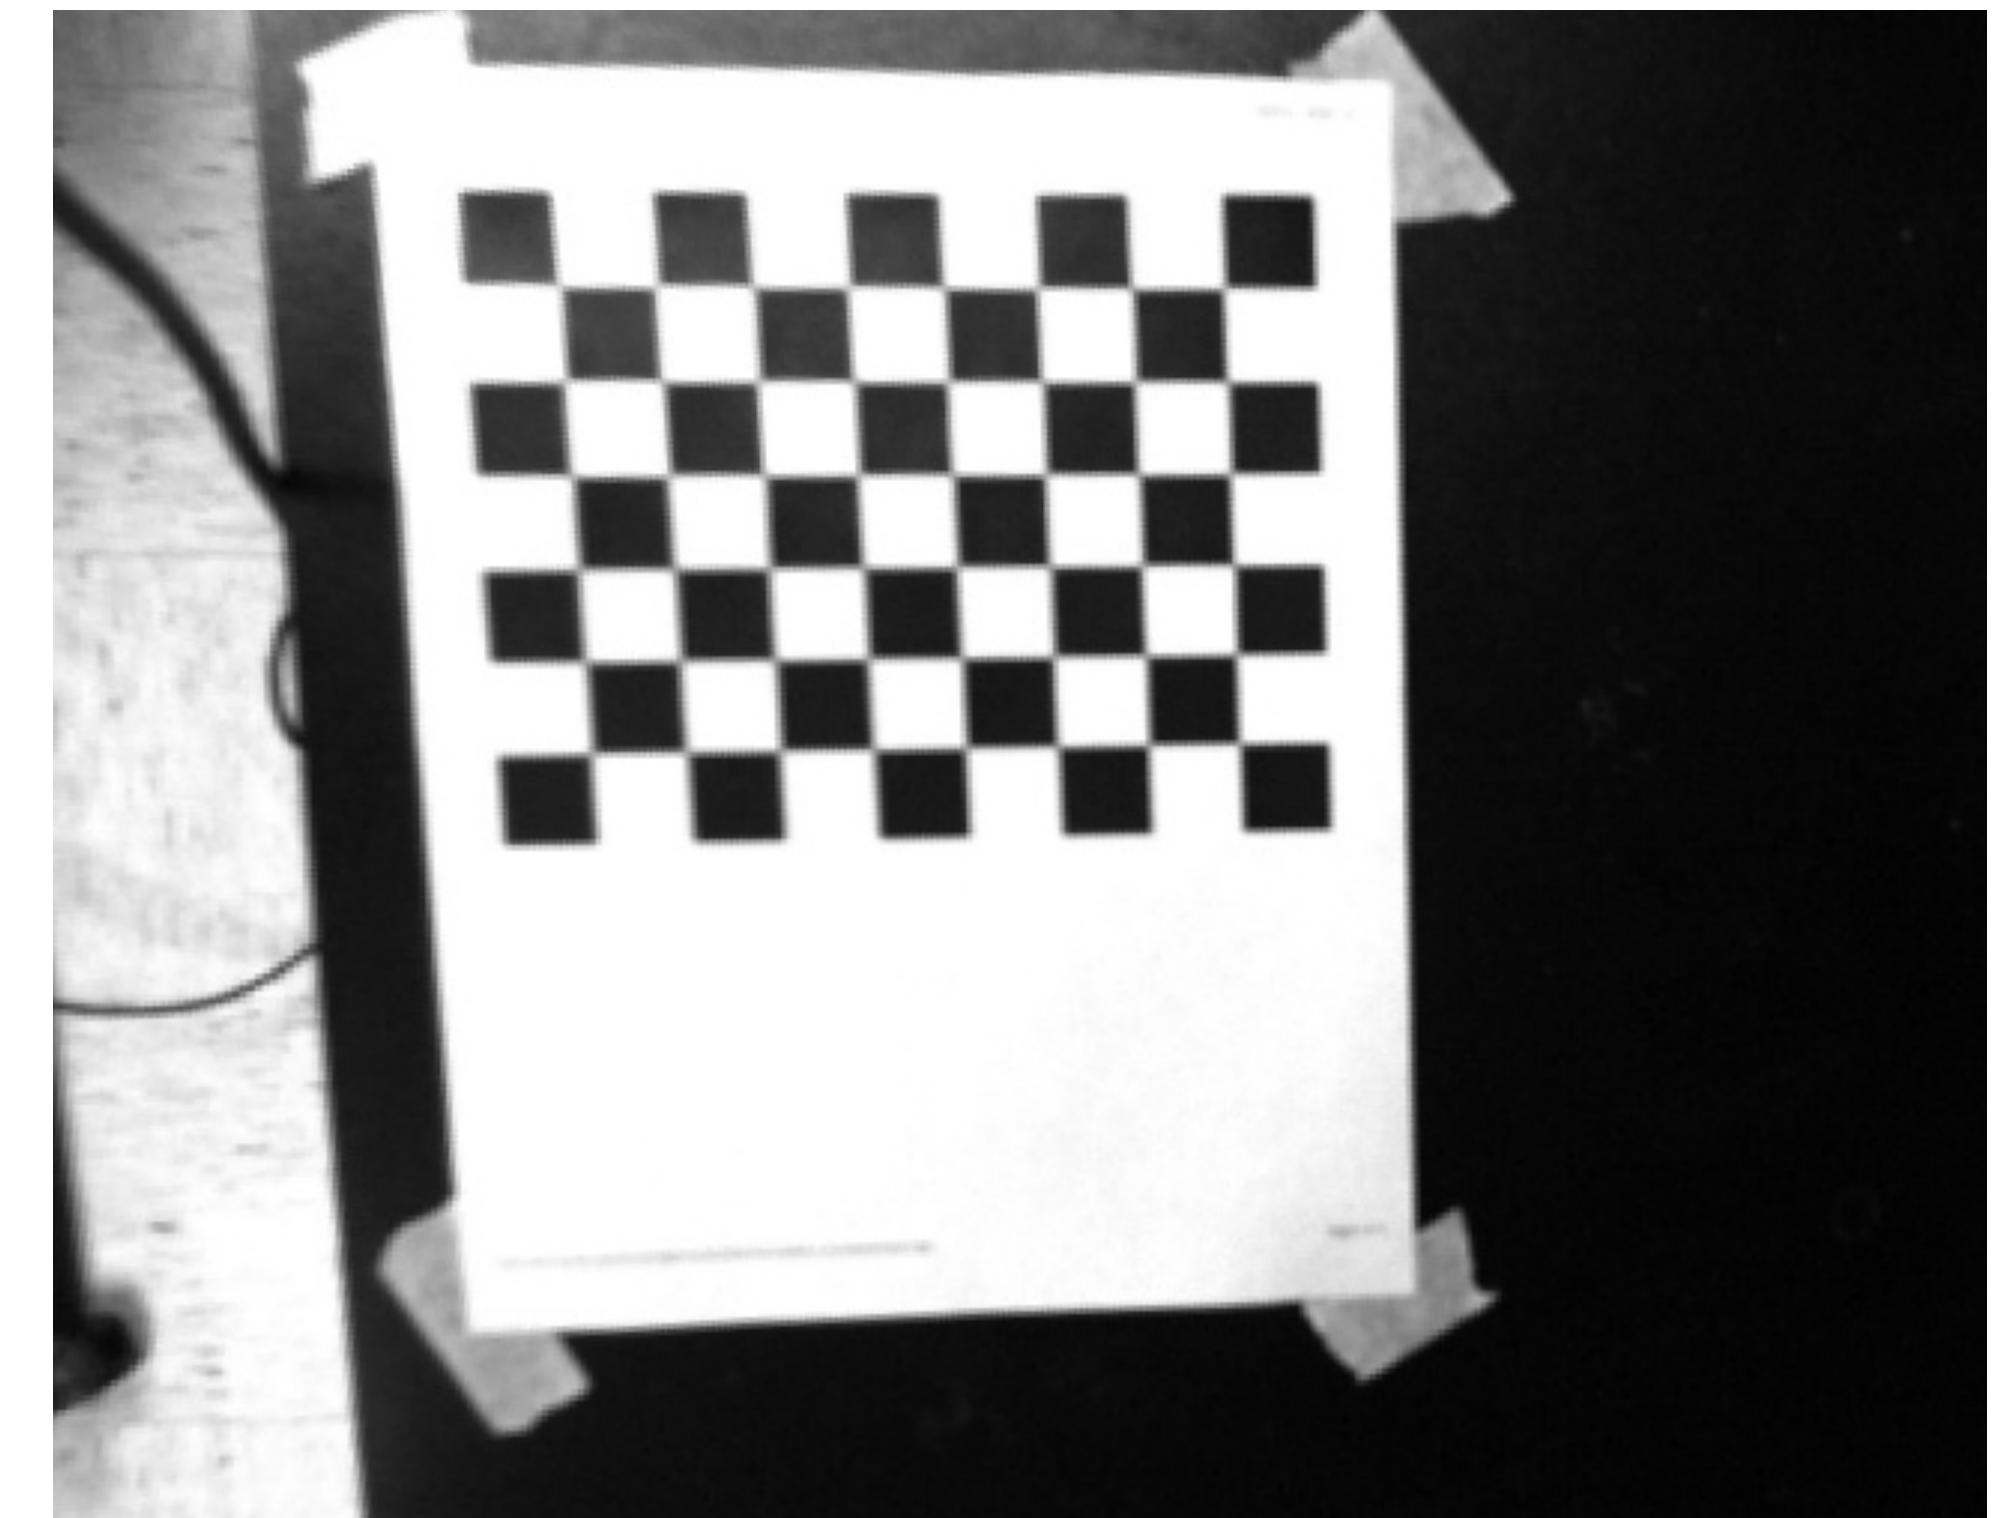
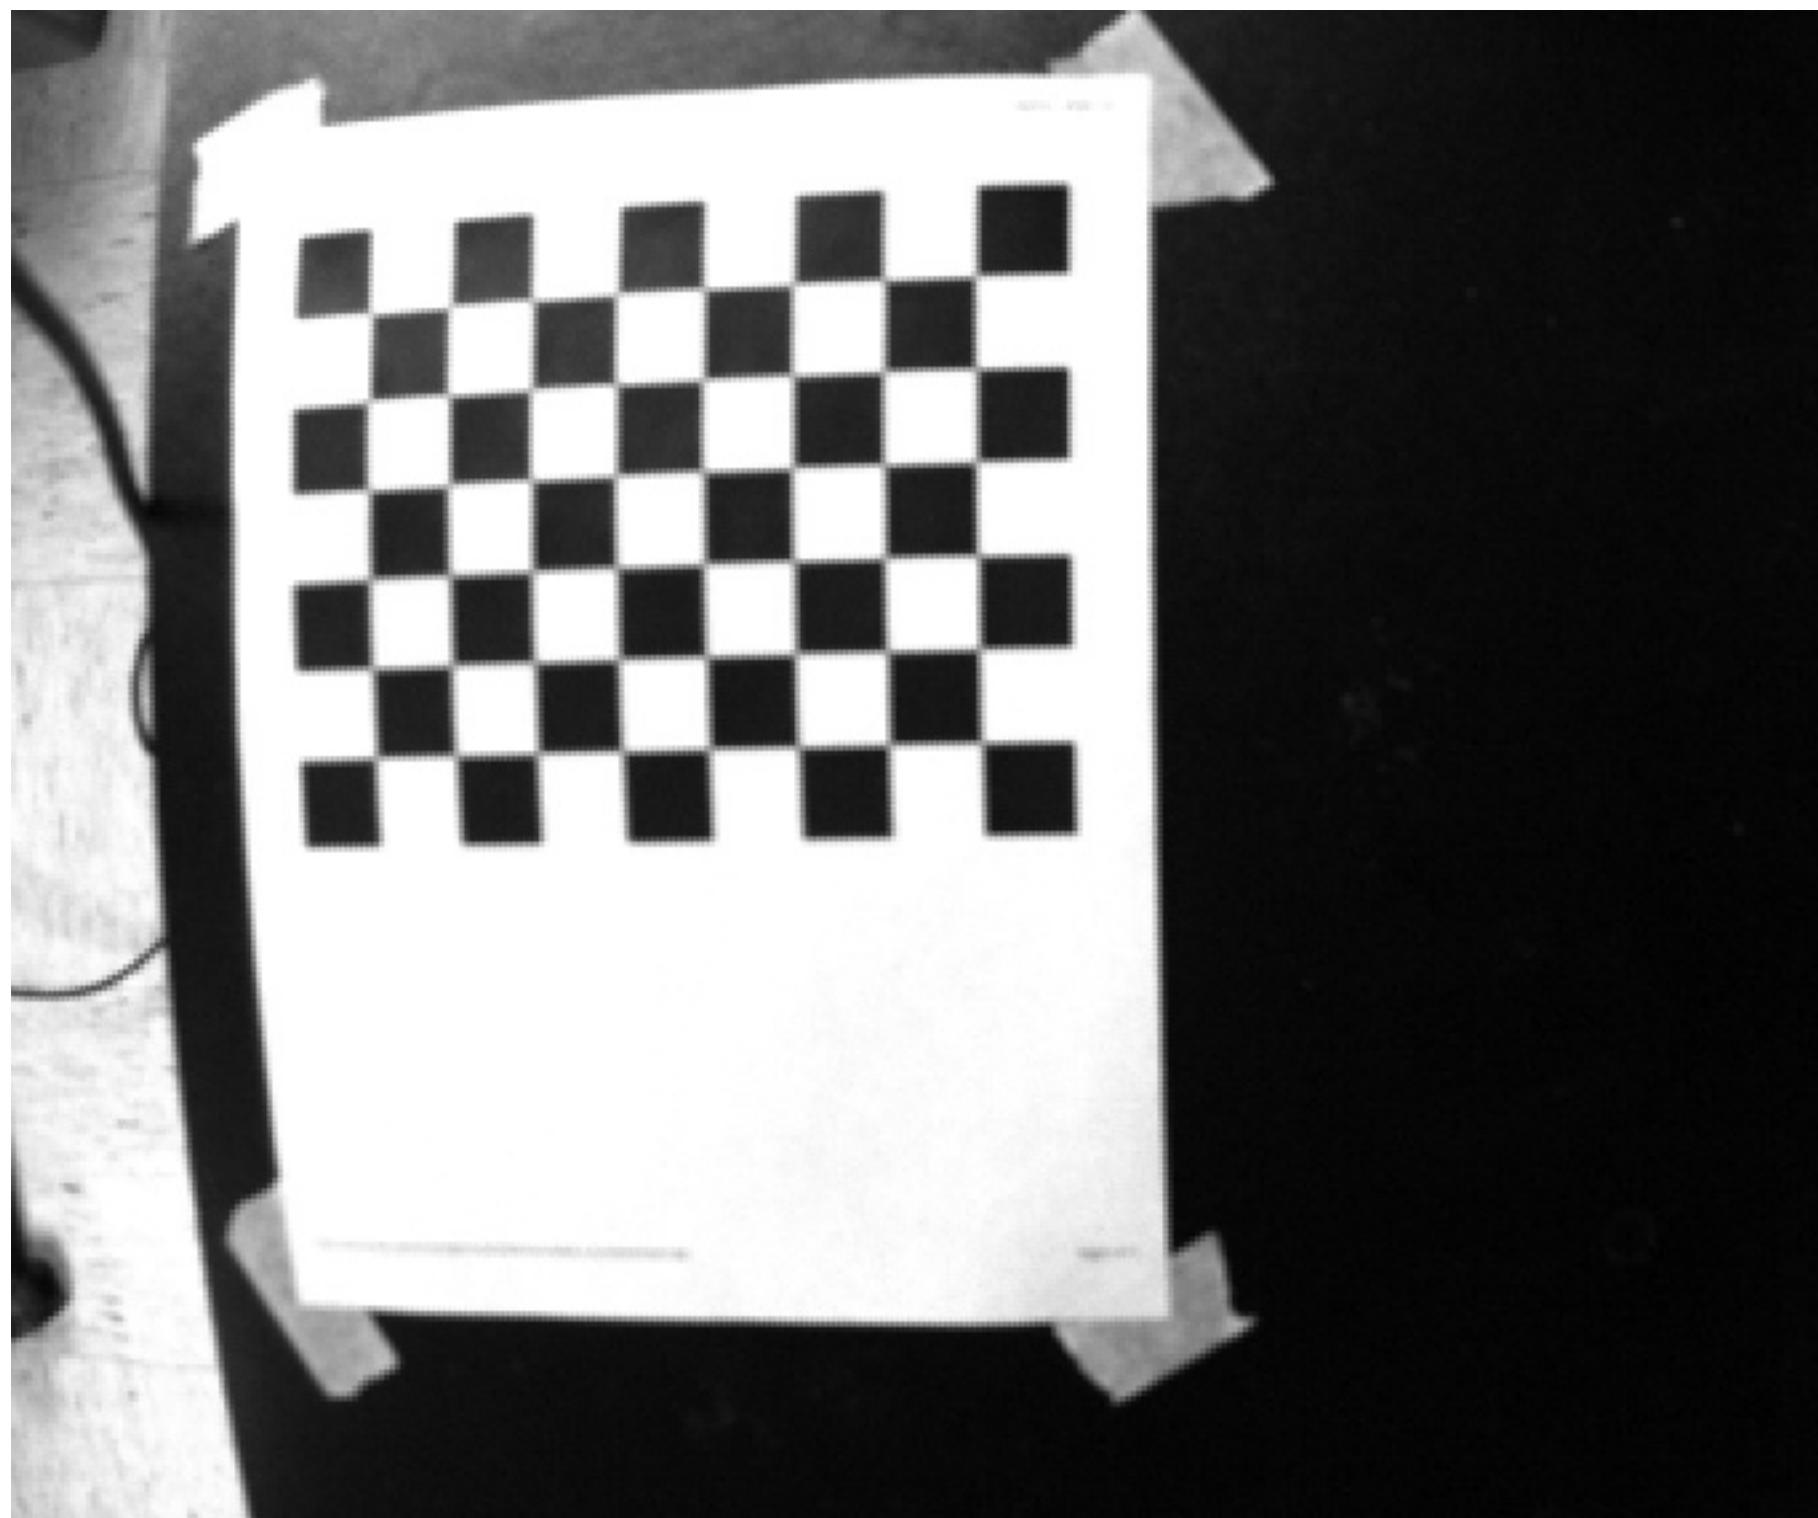
The mouse is very dark,
relative to the background
Thresholding in the value
channel might be sufficient

Hands on: Contours & Basic Object Detection

- Useful contour filtering strategies:
 - Compare to known shape:
 - Circles can be detected by the ratio between perimeter (`cv2.arcLength`) and equivalent diameter $2\sqrt{\text{contourArea} / \pi}$
 - Rectangles and ellipses can be detected by seeing how well a best-fit shape matches the contour (`cv2.fitEllipse`, `cv2.minAreaRect`)
 - Straight lines are more easily located through a Hough transform of the binary image than contour filtering

Uncommon Image Operations

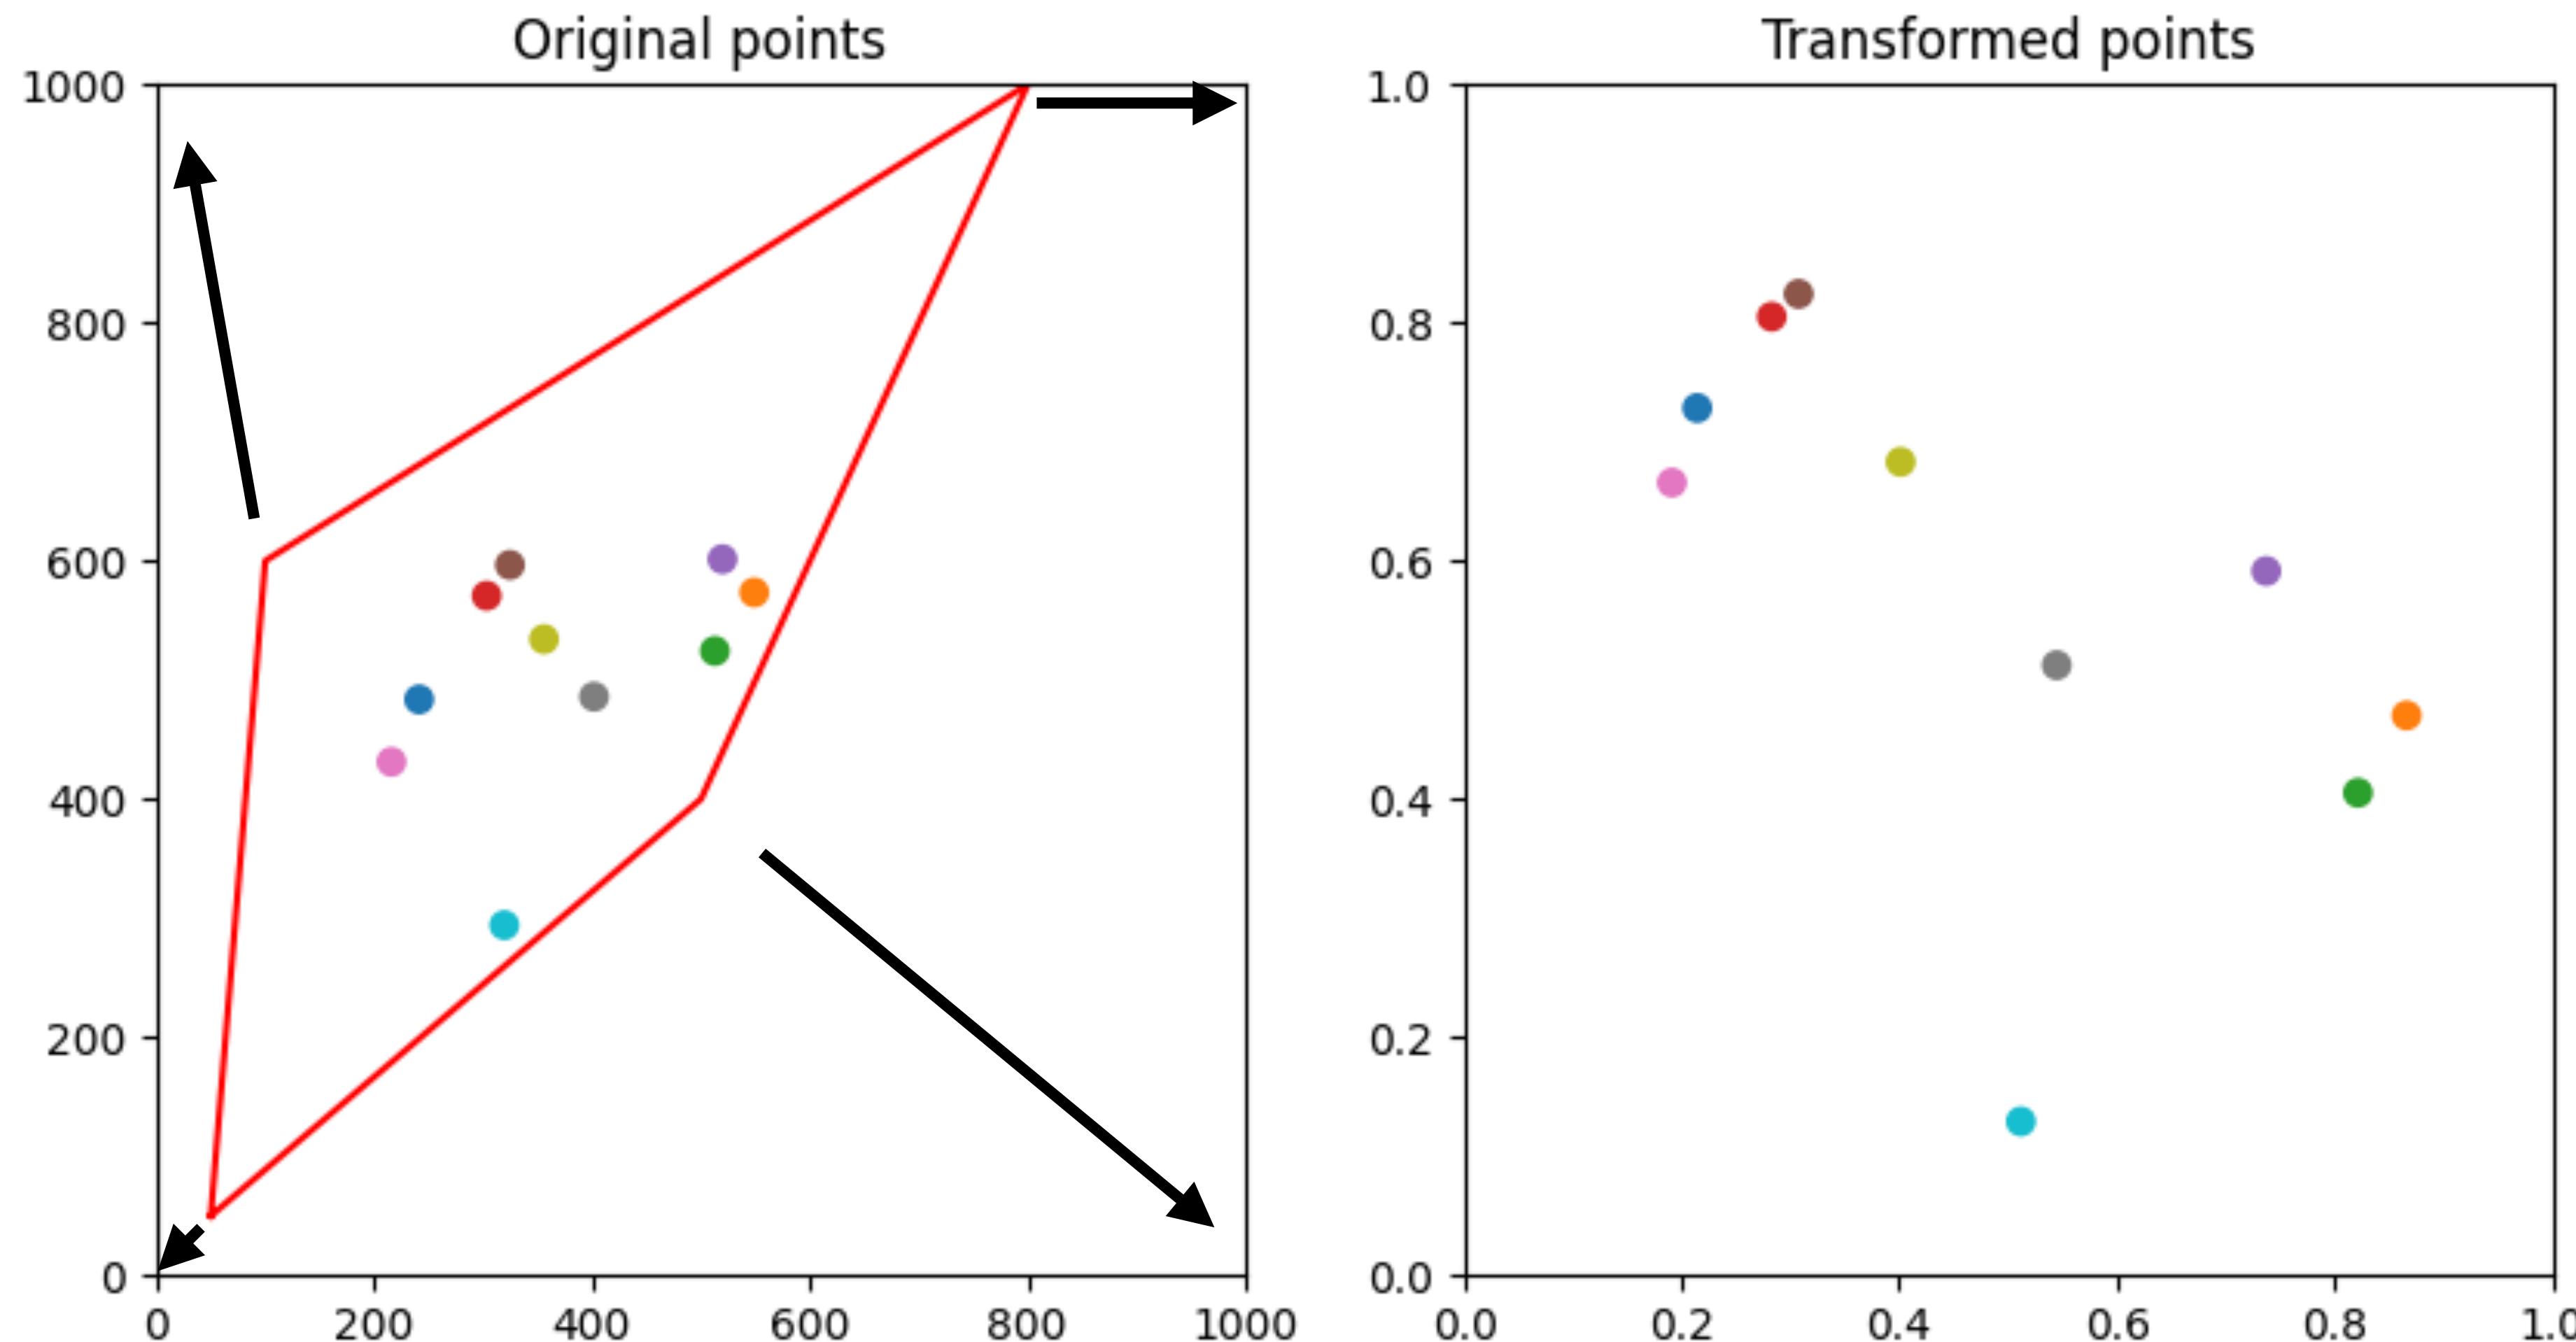
Uncommon Image Operations: Distortion



Uncommon Image Operations: Distortion

- Very situational
- <https://github.com/Aramist/camera-calibration>

Uncommon Image Operations: Perspective



Uncommon Image Operations: Perspective

- Solves for the affine transformation which best maps one quadrilateral to another
- Useful for 2D coordinate frame conversions
- Can warp entire images!
- See [cv2.warpPerspective](#)

```
src_points = np.array([
    [50, 50],
    [500, 400],
    [800, 1000],
    [100, 600],
]).astype(np.float32)

dst_points = np.array([
    [0.0, 0.0],
    [1.0, 0.0],
    [1.0, 1.0],
    [0.0, 1.0],
]).astype(np.float32)

H = cv2.getPerspectiveTransform(src_points, dst_points)

test_points = np.random.uniform(np.array([200, 200]), np.array([700, 700]), (10, 2))
test_points = np.concatenate([test_points, np.ones((10, 1))], axis=1) # add 1 to account for the bias term
transformed_points = np.einsum('ij,bj->bi', H, test_points)
transformed_points = transformed_points[:, :2] / transformed_points[:, 2:3]
```

Working with GUIs

Working with GUIs

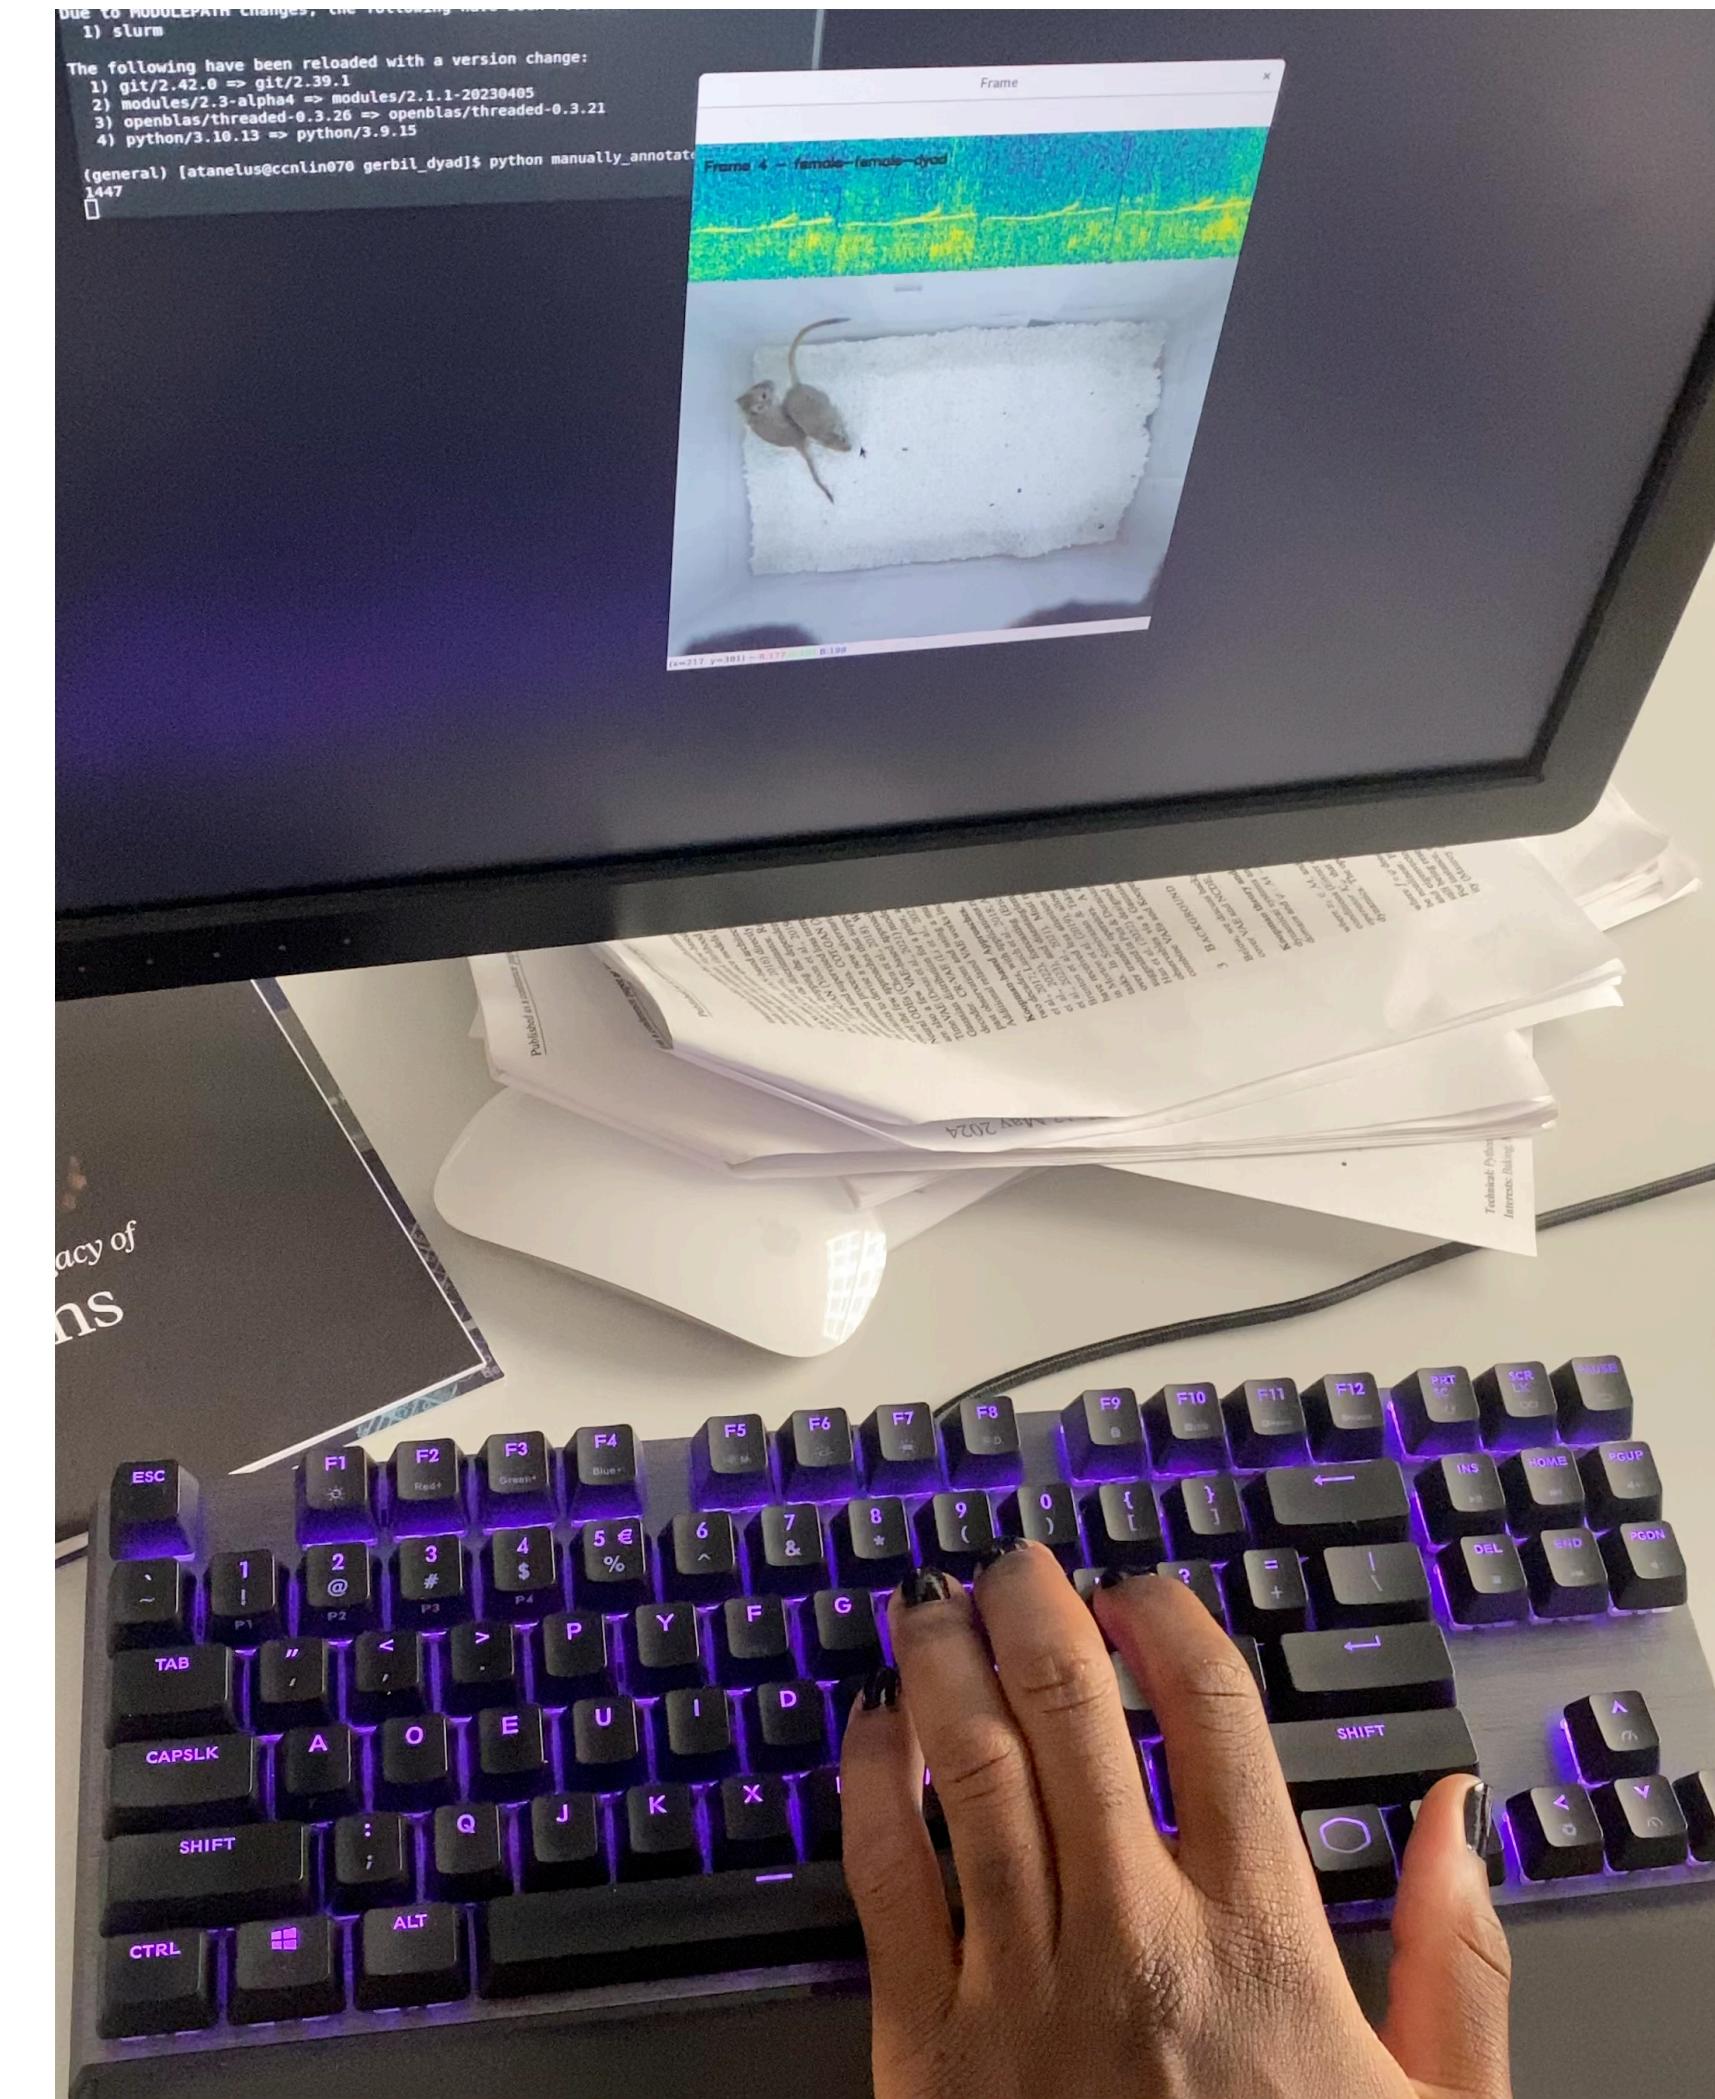
- Create a window: `cv2.namedWindow(name, mode)`
 - I usually use `cv2.WINDOW_NORMAL` (resizable, ratio locked)
- Show an image: `cv2.imshow(windowName, image)`
- IMPORTANT receive input: `cv2.waitKey(max_wait_ms)`
 - Receives keyboard input from the active window
 - Important because calling this prompts opencv to update the window and show any queued images
- Destroy windows: `cv2.destroyAllWindows()`

Working with GUIs: Advanced

- Mouse callbacks: `cv2.setMouseCallback(win_name, callback_fn)`
 - `def callback_fn(event, x, y, flags, param)`
 - Event is an enum for the event type(l click down, r click up, etc.)
 - (x,y) is the position of the mouse in the image displayed on the window
 - I've never touched the other two args

Working with GUIs: Case Study

Gerbil dyadic interaction data annotation in
<300loc : <https://gist.github.com>



Working with Videos

Working with Videos: Reading Videos

- Videos are loaded through the cv2.VideoCapture class
- Initialize the class with a string (path to a video) or an integer (index of the webcam to use)
- Grab frames with read: returns a boolean (whether a frame was successfully grabbed) and an image
- ✨live demo✨ (what could go wrong?)

Working with Videos: Writing Videos

- Videos are written with the cv2.VideoWriter class. This is one of the more inconveniently designed features of OpenCV
- Initialize with cv2.VideoWriter(filename: str, fourcc: int, fps: float, size: tuple[int, int], isColor: bool)
 - Filename does not accept path-like objects, only str
 - If you end up with 512-byte empty videos, double check your size arg. OpenCV does not throw errors if you feed in an image of the wrong size, rather it silently fails to write.
 - If you fail to call release() at the end of your script, you will likewise end up with an empty video (on my machine)

Working with Videos: Writing Videos

- What is a fourcc code?
 - Describes the codec used to encode / compress the video. The codes available to you depend on your system and the extension you wish to save the video as. Passing in a value of -1 into the constructor will prompt OpenCV to list all that are available on your system for a given filetype.
 - I commonly use 'DIVX' for .avi videos and 'mp4v' for .mp4 (capitalization matters)
 - To get the fourcc integer code from a string, use
`cv2.VideoWriter_fourcc(*string_code)`
- ✨live demo✨

Working with Videos: Writing Videos

- Matplotlib figures can be converted to images in OpenCV without much overhead:

```
def get_frame_buffer(figure: mpl.figure.Figure) -> np.ndarray:  
    """Returns a BGR image for opencv from a matplotlib figure"""  
    figure.canvas.draw()  
    buf = figure.canvas.buffer_rgba()  
    width, height = figure.canvas.get_width_height()  
    return np.frombuffer(buf, np.uint8).reshape(height, width, 4)[..., 2::-1]
```

Challenge: Tracking Football Players